


Feature point description and matching

IN4393 – Computer Vision

Introduction

- Last week, we learned how to detect stable, invariant feature points
- How can we use these feature points to match objects in images?



Feature point descriptors

Rotation-invariance

- Like keypoint detectors, descriptors should be scale- and rotation-invariant:
 - However, simple rotation-invariant descriptors have poor discriminability

Rotation-invariance


- Like keypoint detectors, descriptors should be scale- and rotation-invariant:
 - However, simple rotation-invariant descriptors have poor discriminability
- A better approach is to estimate the *dominant orientation* at a keypoint:
 - Simple approach estimates dominant orientation from the Gaussian-weighted horizontal and vertical gradients:

$$\alpha(\mathbf{x}) = \text{atan2} \left(\sum_{\mathbf{x}_i \in \mathcal{N}_{\mathbf{x}}} w(\mathbf{x}_i) I_y(\mathbf{x}_i), \sum_{\mathbf{x}_i \in \mathcal{N}_{\mathbf{x}}} w(\mathbf{x}_i) I_x(\mathbf{x}_i) \right)$$

- Could also look at the principal eigenvector of the second-order matrix


Rotation-invariance

- When the gradients are small, *orientation histograms* are more reliable:




* SIFT uses 36 orientation bins here

- Illustration of dominant-orientation estimation at multiple scales:





SIFT Descriptor

- The SIFT descriptor measures orientation histograms in small blocks:



(a) image gradients



(b) keypoint descriptor

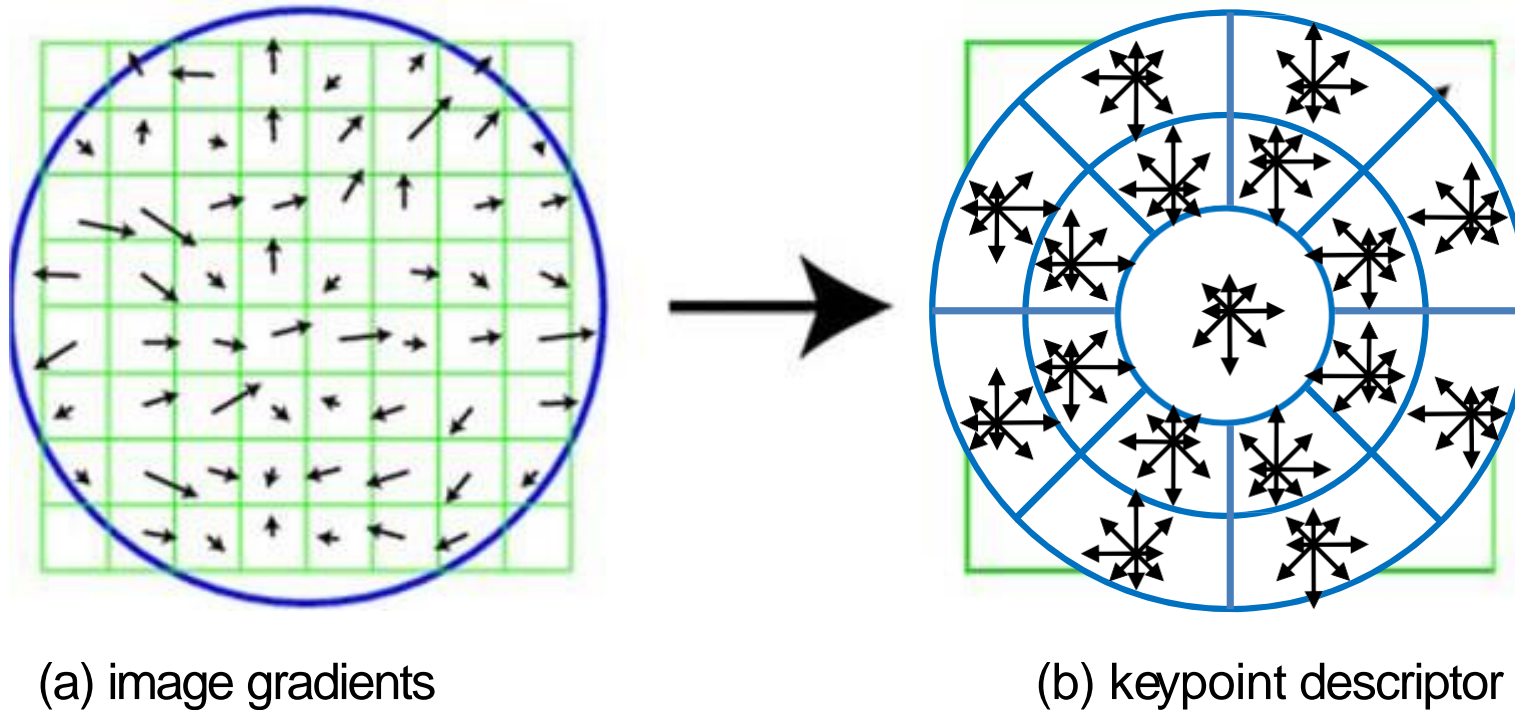
- Image gradients taken *at the right scale, relative to the dominant orientation*

SIFT Descriptor

- Details of the standard SIFT descriptor:
 - Uses a 16x16 pixel window, with 4x4 pixel quadrants and 8 orientation bins
 - Soft addition of gradient magnitudes to histogram bins using trilinear interpolation
 - The 128D descriptor is normalized to unit-length to reduce contrast and gain effects
 - Values are clipped to 0.2 and descriptor is renormalized to deal with high contrast
 - Sometimes PCA is applied to reduce dimensionality of SIFT descriptors (PCA-SIFT)

GLOH descriptor

- Variant of SIFT that uses *log-polar* binning structure:




- Intuition behind GLOH (**Gradient Location and Orientation Histogram**) is that you should not split up the central image part around the feature point

Feature point matching

Feature point matching


- The goal is to find corresponding points in two images for further processing:



- Key idea behind feature point matching:
 - Corresponding points have similar feature point descriptors

Matching feature points


- Similar points have similar descriptors, *i.e.* are nearby in “*descriptor space*”:



- Simple approach to finding matches is to threshold Euclidean distances
 - It may be necessary to *whiten* the descriptors first: normalize feature “scales”

Matching feature points


- Similar points have similar descriptors, *i.e.* are nearby in “*descriptor space*”:



- Simple approach to finding matches is to threshold Euclidean distances
 - It may be necessary to *whiten* the descriptors first: normalize feature “scales”


Matching feature points

- Setting threshold on Euclidean distance is hard; optimal value may vary a lot
- Moreover, a single feature point may get many potential matches
- A better approach is to restrict matches to *nearest neighbors* only:




Matching feature points

- Setting threshold on Euclidean distance is hard; optimal value may vary a lot
- Moreover, a single feature point may get many potential matches
- A better approach is to restrict matches to *nearest neighbors* only:




Matching feature points

- *Nearest neighbor distance ratio:* $NNDR = \frac{d(\text{target}, \text{nearest neighbor 1})}{d(\text{target}, \text{nearest neighbor 2})}$
- Compares the distance to the first neighbor with that to the second neighbor:




Matching feature points

- *Nearest neighbor distance ratio:* $NNDR = \frac{d(\text{target}, \text{nearest neighbor 1})}{d(\text{target}, \text{nearest neighbor 2})}$
- Compares the distance to the first neighbor with that to the second neighbor:



Evaluating quality of matching

- *True positives* (TP): number of matches that were correctly detected
- *False positives* (FP): number of non-matches that were erroneously detected
- *True negatives* (TN): number of non-matches that were correctly rejected
- *False negatives* (FN): number of matches that were not detected




Evaluating quality of matching

- *True positive rate*: percentage of true matches that is indeed proposed
- *False positive rate*: percentage of non-matches that is erroneously proposed

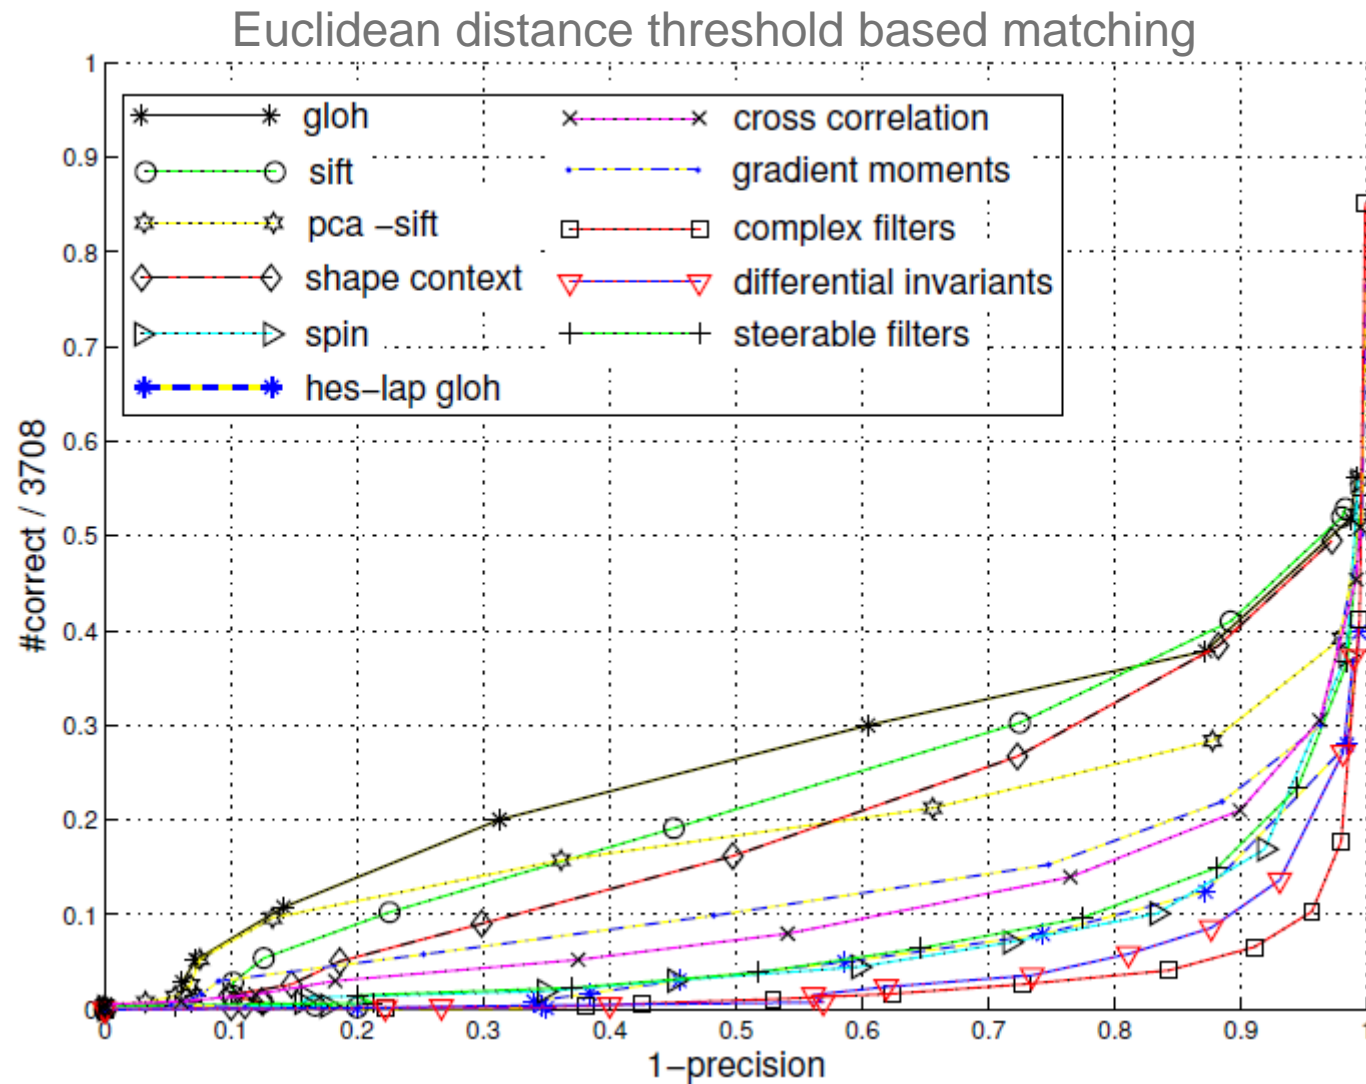
$$TPR = \frac{TP}{TP + FN} = \frac{TP}{P} \quad FPR = \frac{FP}{FP + TN} = \frac{FP}{N}$$

- *Receiver-operating characteristic (ROC) curve* relates TPR and FPR:



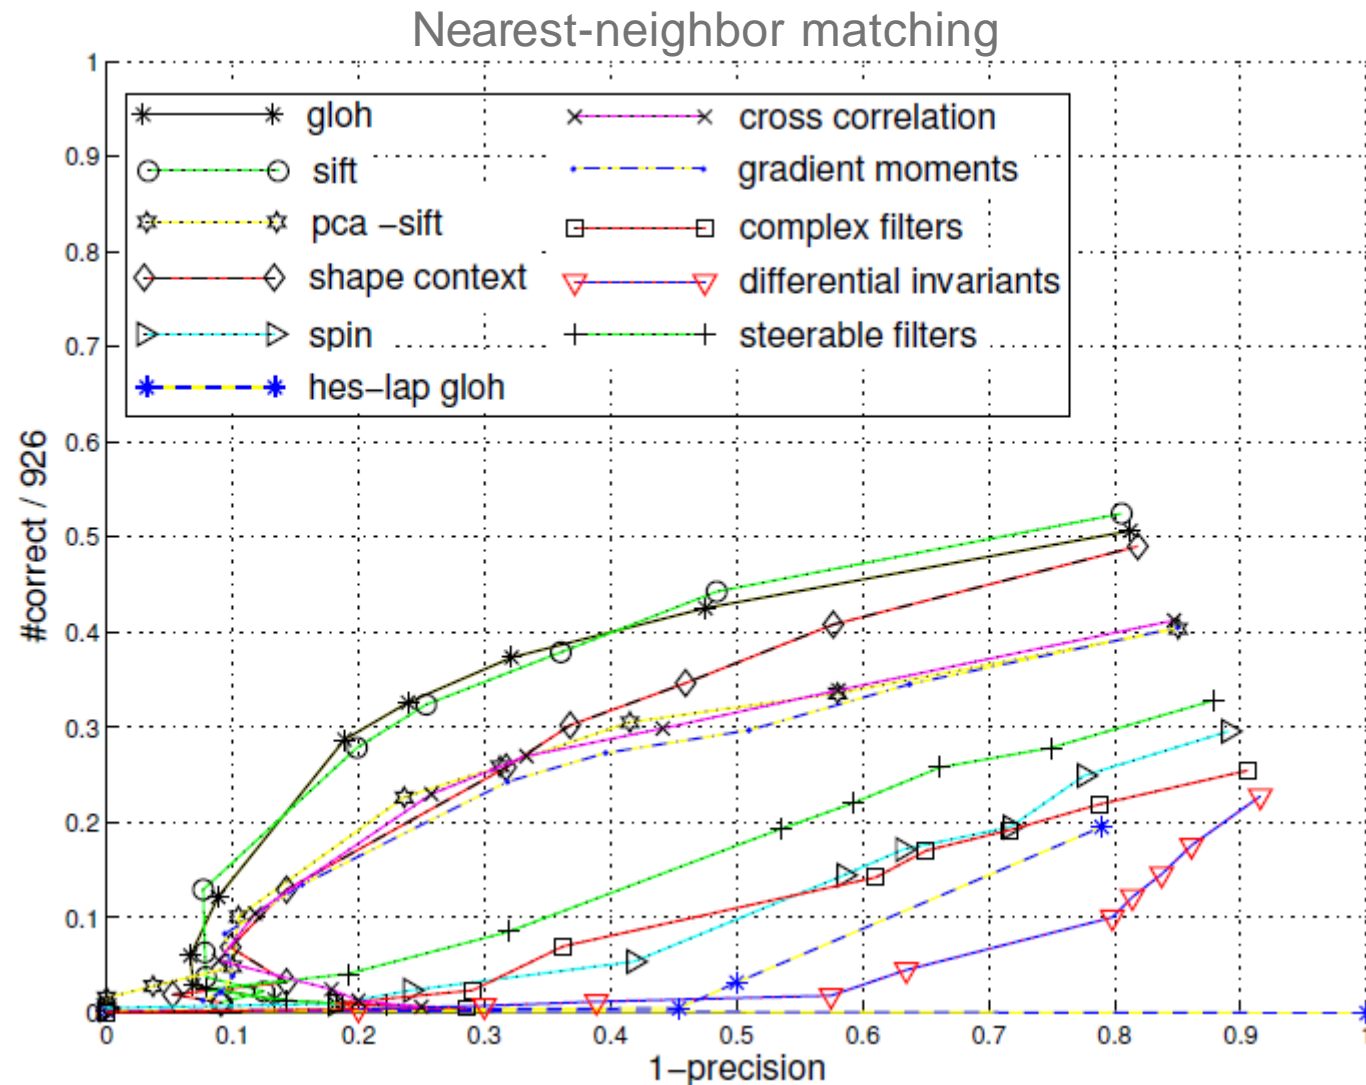
Comparing matching approaches

- Euclidean vs. nearest-neighbor vs. NNDR for various descriptors:



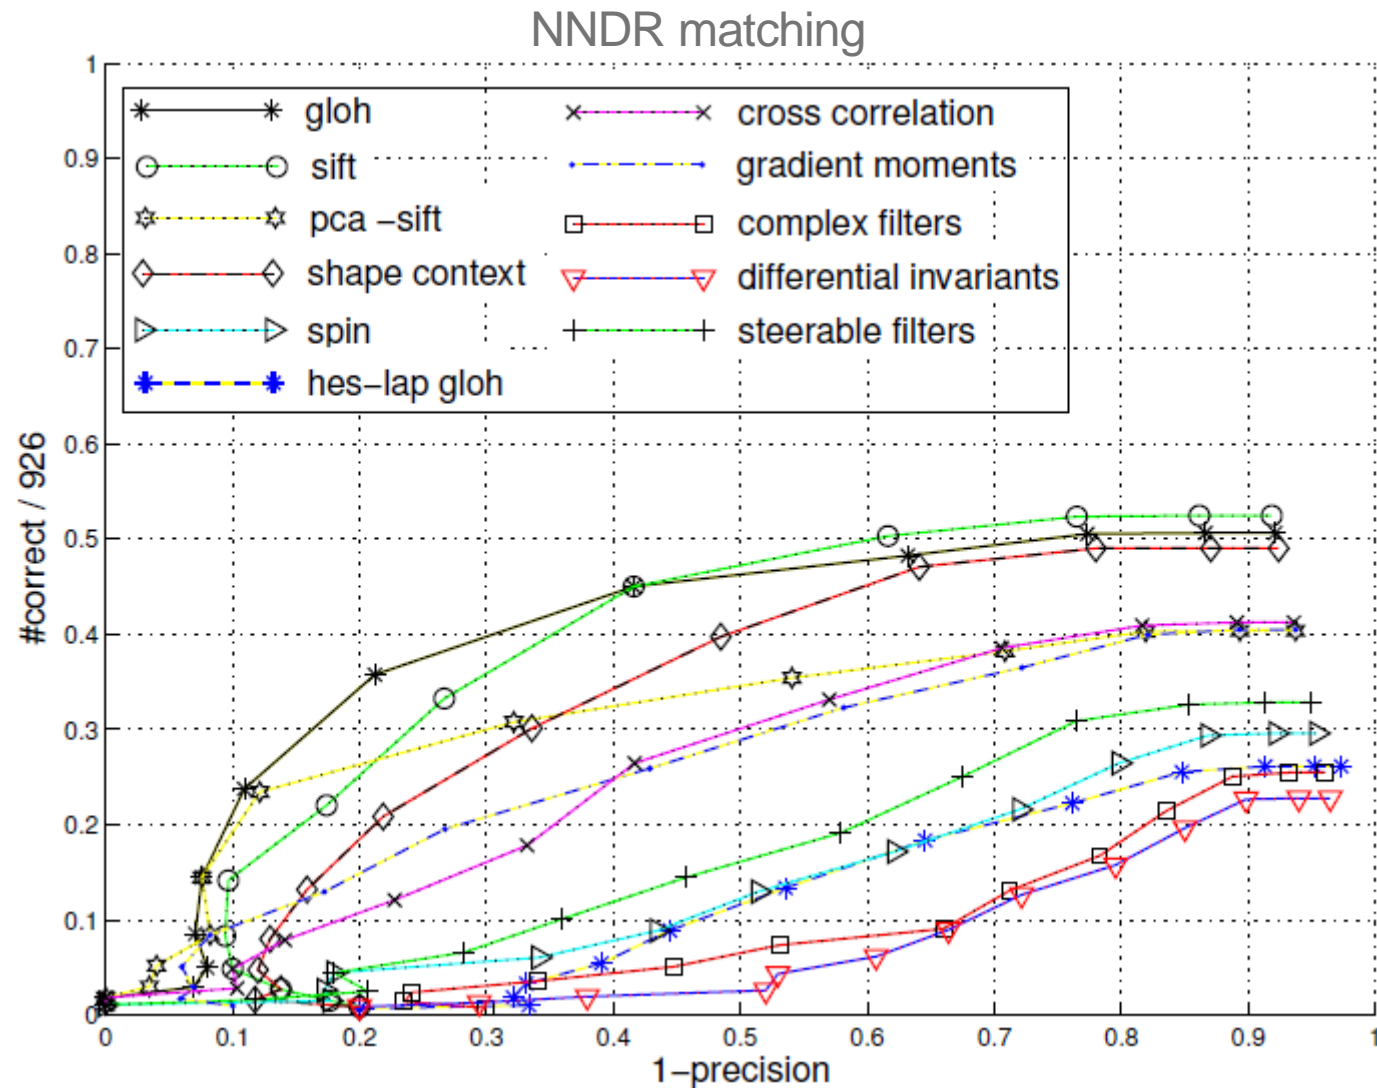
Comparing matching approaches

- Euclidean vs. nearest-neighbor vs. NNDR for various descriptors:




Comparing matching approaches

- Euclidean vs. nearest-neighbor vs. NNDR for various descriptors:



Example: Object recognition

- Given a database of labeled objects, you could do *object recognition*:






Example: Object recognition

- Given a database of labeled objects, you could do *object recognition*:




Efficient matching

- Naively matching two sets of feature points is $\mathcal{O}(NM)$
- *Locality sensitive hashing* or *kd-trees* may speed up nearest neighbor search




Locality-sensitive hashing

- LSH uses hashing functions that take “location” of object in consideration:




Locality-sensitive hashing

- LSH uses hashing functions that take “location” of object in consideration:



- Example of a locality-sensitive hashing function for points in a space:
 - Project the point onto a random subspace; divide result into 4 buckets (2 bits)



Locality-sensitive hashing


- LSH uses hashing functions that take “location” of object in consideration:

- Simple example


- Projection

Whilst normal hash functions try to **minimize** the probability of collision, LSH hash functions try to **maximize** probability of similar items colliding.

(2 bits)




Locality-sensitive hashing



- Mathematically, we could express this locality sensitive hash function as:

$$h(\mathbf{x}) = \begin{cases} 0 & \text{if } \mathbf{w}^\top \mathbf{x} \leq -\tau \\ 1 & \text{if } -\tau < \mathbf{w}^\top \mathbf{x} \leq 0 \\ 2 & \text{if } 0 < \mathbf{w}^\top \mathbf{x} \leq \tau \\ 3 & \text{if } \mathbf{w}^\top \mathbf{x} > \tau \end{cases}$$

Locality-sensitive hashing




- Mathematically, we could express this locality sensitive hash function as:

$$h(\mathbf{x}) = \begin{cases} 0 & \text{if } \mathbf{w}^\top \mathbf{x} \leq -\tau \\ 1 & \text{if } -\tau < \mathbf{w}^\top \mathbf{x} \leq 0 \\ 2 & \text{if } 0 < \mathbf{w}^\top \mathbf{x} \leq \tau \\ 3 & \text{if } \mathbf{w}^\top \mathbf{x} > \tau \end{cases}$$

**random
projection**

Locality-sensitive hashing



- Mathematically, we could express this locality sensitive hash function as:


$$h(\mathbf{x}) = \begin{cases} 0 & \text{if } \mathbf{w}^\top \mathbf{x} \leq -\tau \\ 1 & \text{if } -\tau < \mathbf{w}^\top \mathbf{x} \leq 0 \\ 2 & \text{if } 0 < \mathbf{w}^\top \mathbf{x} \leq \tau \\ 3 & \text{if } \mathbf{w}^\top \mathbf{x} > \tau \end{cases}$$

random
projection

threshold
parameter


Locality-sensitive hashing

- Retrieval of nearest neighbors of a query point q using LSH works as follows:




Locality-sensitive hashing

- Retrieval of nearest neighbors of a query point q using LSH works as follows:
 - Hash all data points using locality-sensitive hash




Locality-sensitive hashing

- Retrieval of nearest neighbors of a query point q using LSH works as follows:
 - Hash all data points using locality-sensitive hash
 - Compute locality-sensitive hash of query point




Locality-sensitive hashing

- Retrieval of nearest neighbors of a query point q using LSH works as follows:
 - Hash all data points using locality-sensitive hash
 - Compute locality-sensitive hash of query point
 - All data points in the bucket are candidate near neighbors



Locality-sensitive hashing

- Retrieval of nearest neighbors of a query point q using LSH works as follows:
 - Hash all data points using locality-sensitive hash
 - Compute locality-sensitive hash of query point
 - All data points in the bucket are candidate near neighbors
 - Compute distances to candidate points to find true nearest neighbors




Locality-sensitive hashing

- LSH projections may be “unlucky” in two main ways:

Locality-sensitive hashing


- LSH projections may be “unlucky” in two main ways:




“Collision”: Distant points
hashed in the same bucket

Locality-sensitive hashing

- LSH projections may be “unlucky” in two main ways:




“Collision”: Distant points hashed in the same bucket



“Split”: Nearby points hashed in different buckets


Locality-sensitive hashing

- Using multiple projections in an LSH resolves “collisions”:



Locality-sensitive hashing


- Using multiple projections in an LSH resolves “collisions”:



- The LSH is given by a concatenation of all individual buckets


Locality-sensitive hashing

- Using multiple separate hash tables when doing LSH resolves “splits”:



Locality-sensitive hashing

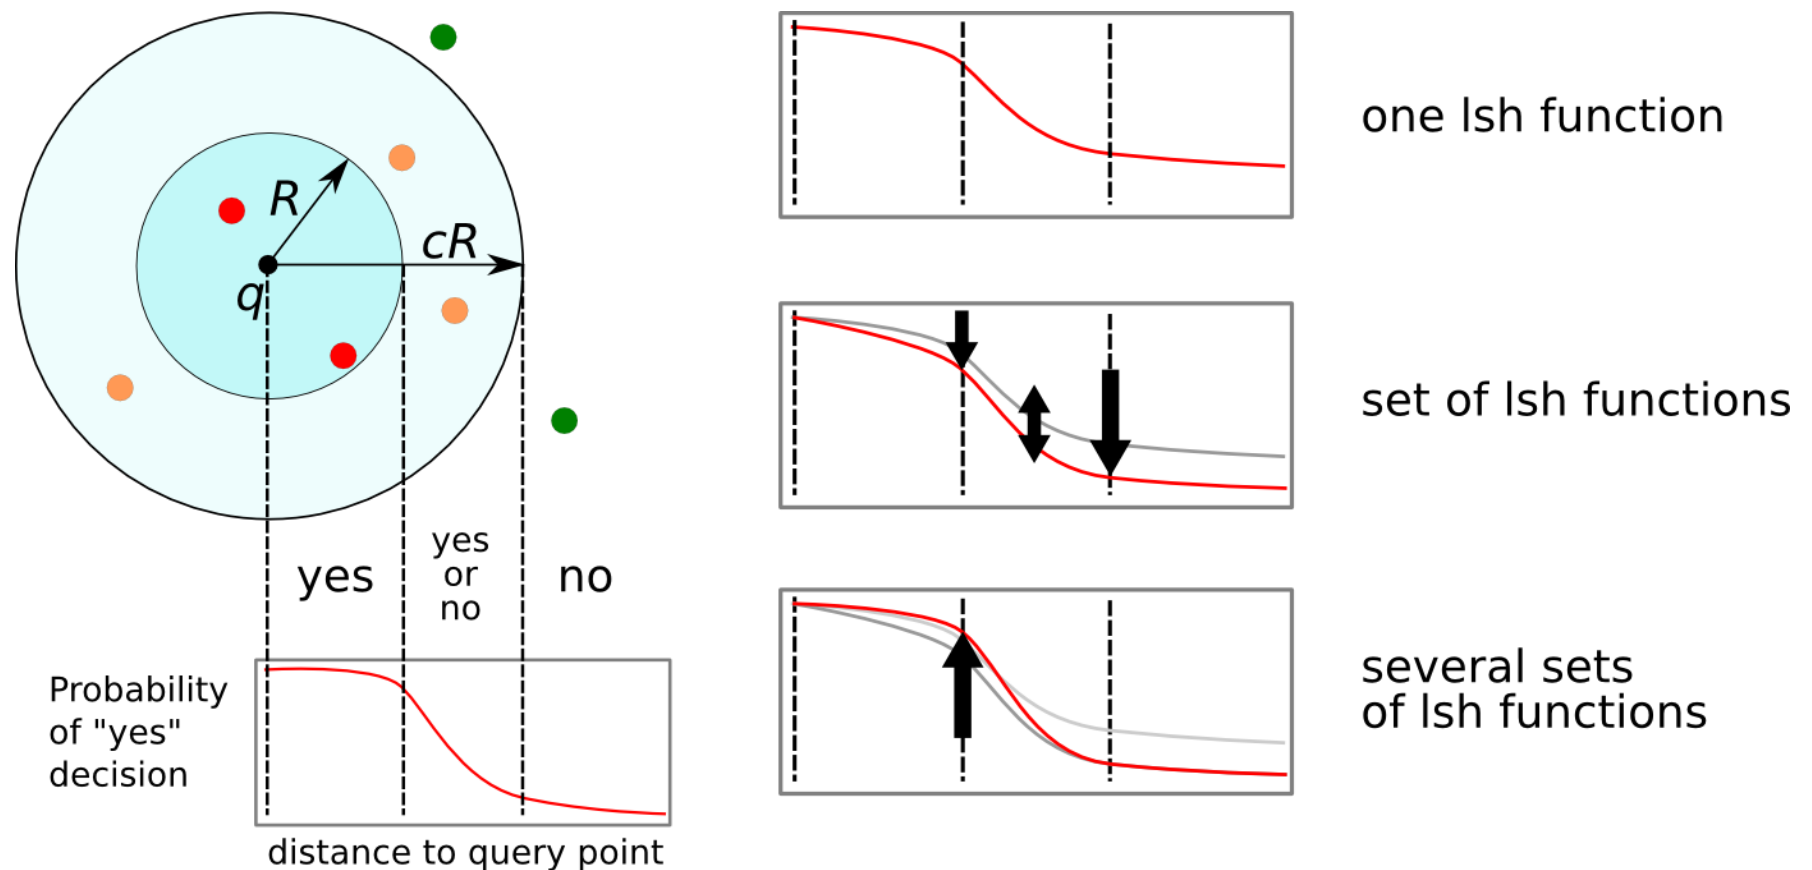
- Using multiple separate hash tables when doing LSH resolves “splits”:



- Points are candidate neighbors if candidate in *any* of the hash tables



Locality-sensitive hashing

- Error analysis of locality-sensitive hashing:



kd-trees

- Tree structure that optimally splits a random dimension at each level:




kd-trees

- Finding the nearest neighbor of a target point using a kd-tree:
 - Identify the bin in which the target point is located
 - Compute distance to nearest neighbor inside this bin (D_{cur})
 - Perform depth-first search on the kd-tree:
 - Prune all cells that are further away than D_{cur}
 - If we arrive at a leaf, search for nearer neighbors, and update D_{cur}

SIFT Matches


- Matches SIFT finds are not generally free of errors:



Feature-based alignment

Feature-based alignment

- Estimating the motion between two images based on set of matched points
- Basic collection of 2D (planar) *coordinate transformations*:



Feature-based alignment

- Basic 2D coordinate transformations $\mathbf{x}' = f(\mathbf{x}; \mathbf{p})$:

Transform	Matrix	Parameters p	Jacobian J
translation	$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$	(t_x, t_y)	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Euclidean	$\begin{bmatrix} c_\theta & -s_\theta & t_x \\ s_\theta & c_\theta & t_y \end{bmatrix}$	(t_x, t_y, θ)	$\begin{bmatrix} 1 & 0 & -s_\theta x - c_\theta y \\ 0 & 1 & c_\theta x - s_\theta y \end{bmatrix}$
similarity	$\begin{bmatrix} 1+a & -b & t_x \\ b & 1+a & t_y \end{bmatrix}$	(t_x, t_y, a, b)	$\begin{bmatrix} 1 & 0 & x & -y \\ 0 & 1 & y & x \end{bmatrix}$
affine	$\begin{bmatrix} 1+a_{00} & a_{01} & t_x \\ a_{10} & 1+a_{11} & t_y \end{bmatrix}$	$(t_x, t_y, a_{00}, a_{01}, a_{10}, a_{11})$	$\begin{bmatrix} 1 & 0 & x & y & 0 & 0 \\ 0 & 1 & 0 & 0 & x & y \end{bmatrix}$

- These transformations use *augmented vector* representation: $\mathbf{x} = [x \ y \ 1]^T$

Feature-based alignment

- *Least squares* provides a simple way to estimate parameters **p** of transform
- Assuming a set of matched feature points $\{(\mathbf{x}_i, \mathbf{x}'_i)\}_{i=1,\dots,N}$, we minimize:


$$E = \sum_i \|f(\mathbf{x}_i; \mathbf{p}) - \mathbf{x}'_i\|^2$$

- If transformation is linear in parameters, closed-form solution exists!


Linear least squares

- Consider the *linear least squares* optimization problem:

$$g(\mathbf{x}) = \|\mathbf{Ax} - \mathbf{b}\|^2$$


residual r (linear in parameters)


$$\mathbf{x}^* = \operatorname{argmin}_{\mathbf{x}} g(\mathbf{x})$$


optimal parameters


Linear least squares

- Consider the *linear least squares* optimization problem:

$$g(\mathbf{x}) = \|\mathbf{Ax} - \mathbf{b}\|^2$$


residual r (linear in parameters)

$$\mathbf{x}^* = \underset{\mathbf{x}}{\operatorname{argmin}} g(\mathbf{x})$$


optimal parameters

- We can solve such an optimization problem by *setting the gradient to zero*:

$$\frac{\partial g}{\partial \mathbf{x}} = 2\mathbf{A}^T(\mathbf{Ax} - \mathbf{b}) = 0$$

$$\mathbf{A}^T \mathbf{Ax} = \mathbf{A}^T \mathbf{b}$$

$$\mathbf{x} = \underbrace{(\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}}$$


pseudo-inverse of \mathbf{A}

Feature-based alignment

- For translation, similarity and affine transforms, the movement is *linear* in \mathbf{p} :

$$\begin{aligned} E &= \sum_i \|f(\mathbf{x}_i; \mathbf{p}) - \mathbf{x}'_i\|^2 = \sum_i \|\mathbf{x}_i + J(\mathbf{x}_i)\mathbf{p} - \mathbf{x}'_i\|^2 \\ &= \sum_i \|J(\mathbf{x}_i)\mathbf{p} - \Delta\mathbf{x}_i\|^2 \end{aligned}$$


with: $\Delta\mathbf{x}_i = \mathbf{x}'_i - \mathbf{x}_i$

- The optimal solution for the problem is thus given in closed form:

$$\mathbf{p}^* = \left(\sum_i J(\mathbf{x}_i)^\top J(\mathbf{x}_i) \right)^{-1} \sum_i J^\top(\mathbf{x}_i) \Delta\mathbf{x}_i$$

Example: Panography

- Simple example of feature-based alignment: Panography (translation model)



- Let's work out the optimal translation based on correspondences

Example: Panography

- Our aim is to find the optimal translation based on keypoint correspondences:

$$\begin{aligned} g(t_x, t_y) &= \sum_{n=1}^N \left\| \begin{bmatrix} x_n & y_n & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ t_x & t_y \end{bmatrix} - \begin{bmatrix} x'_n \\ y'_n \end{bmatrix} \right\|^2 \\ &= \sum_{n=1}^N \left\| \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} t_x \\ t_y \end{bmatrix} - \begin{bmatrix} x_n - x'_n \\ y_n - y'_n \end{bmatrix} \right\|^2 \\ &= \left\| \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ \dots & \dots \\ 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} t_x \\ t_y \end{bmatrix} - \begin{bmatrix} x_1 - x'_1 \\ y_1 - y'_1 \\ \dots \\ x_N - x'_N \\ y_N - y'_N \end{bmatrix} \right\|^2 \end{aligned}$$

Example: Panography

- Our aim is to find the optimal translation based on keypoint correspondences:

$$\begin{aligned} g(t_x, t_y) &= \sum_{n=1}^N \left\| \begin{bmatrix} x_n & y_n & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ t_x & t_y \end{bmatrix} - \begin{bmatrix} x'_n \\ y'_n \end{bmatrix} \right\|^2 \\ &= \sum_{n=1}^N \left\| \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} t_x \\ t_y \end{bmatrix} - \begin{bmatrix} x_n - x'_n \\ y_n - y'_n \end{bmatrix} \right\|^2 \\ &= \left\| \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ \dots & \dots \\ 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} t_x \\ t_y \end{bmatrix} - \begin{bmatrix} x_1 - x'_1 \\ y_1 - y'_1 \\ \dots \\ x_N - x'_N \\ y_N - y'_N \end{bmatrix} \right\|^2 \end{aligned}$$

- This can be recognized as a standard *linear least-squares* problem!

RANSAC


- Erroneous matches may highly influence our estimate for t_x and t_y

RANSAC

- Erroneous matches may highly influence our estimate for t_x and t_y
- A *RANSAC* algorithm for fitting a panography would roughly work as follows:
 - *Fit model* on current inliers: solve linear-least squares problem
 - *Determine inliers*: e.g., find points for which $\|x_n + t_x - x'_n\|^2 + \|y_n + t_y - y'_n\|^2$ is small, and consider those as the new inliers
 - Return to the first step using the new inliers
 - Keep track of the best model so far (in terms of the squared error) that has “sufficient” inliers

RANSAC

- *R*andom *S*ample *C*onsensus aims to fit model whilst being robust to outliers:




RANSAC

- RANSAC is an iterative algorithm that works using the following steps:
 - 1) Model is fitted to the *hypothetical inliers*
 - 2) Data are tested against the fitted model to determine hypothetical inliers
 - 3) Return to step 1) until sufficient points are classified as inliers (or fixed number of times)
 - 4) Keep track of best model so far during iterations

RANSAC

- RANSAC is an iterative algorithm that works using the following steps:
 - 1) Model is fitted to the *hypothetical inliers*
 - 2) Data are tested against the fitted model to determine hypothetical inliers
 - 3) Return to step 1) until sufficient points are classified as inliers (or fixed number of times)
 - 4) Keep track of best model so far during iterations
- No upper bound on number of iterations available; RANSAC may take forever
- RANSAC has two magic parameters: How far can data be from the model to be considered inlier? And how much data is needed to accept the model?



Reading material: Section 4 and 6