

Image stitching

IN4393 – Computer Vision

Image stitching

- We are given a bunch of photographs; how do we *stitch* them together?



- But first, we need to understand *non-linear least squares* and *homographies*


Non-linear least squares problems

Non-linear least squares

- Fitting a panography was *linear* in the transformation parameters:

$$\begin{aligned} E &= \sum_i \|f(\mathbf{x}_i; \mathbf{p}) - \mathbf{x}'_i\|^2 = \sum_i \|\mathbf{x}_i + J(\mathbf{x}_i)\mathbf{p} - \mathbf{x}'_i\|^2 \\ &= \sum_i \|J(\mathbf{x}_i)\mathbf{p} - \Delta\mathbf{x}_i\|^2 \end{aligned}$$

- For more complex motion models, the transformation is *non-linear*:






Non-linear least squares

- Consider the *non-linear least squares problem*:

$$g(\mathbf{x}) = \|f(\mathbf{x}; \mathbf{A}) - \mathbf{b}\|^2$$

- This problem is in general not *convex*; it may have multiple *local minima*:




Taylor expansion

- The *Taylor expansion* of the function $f(x)$ around a is given by:

$$\sum_{n=0}^{\infty} \frac{f^{(n)}(a)}{n!} (x - a)^n$$

- Herein, $f^{(n)}$ denotes the n -th derivative




Gauss-Newton method


- Iteratively find parameter updates $\Delta \mathbf{x}$


Gauss-Newton method

- Iteratively find parameter updates $\Delta \mathbf{x}$
- Perform a *first-order Taylor expansion* of the residual around the current \mathbf{x} :

$$\|f(\mathbf{x} - \Delta \mathbf{x}; \mathbf{A}) - \mathbf{b}\|^2 \approx \|f(\mathbf{x}; \mathbf{A}) + J(\mathbf{x})\Delta \mathbf{x} - \mathbf{b}\|^2$$

 parameter update
objective

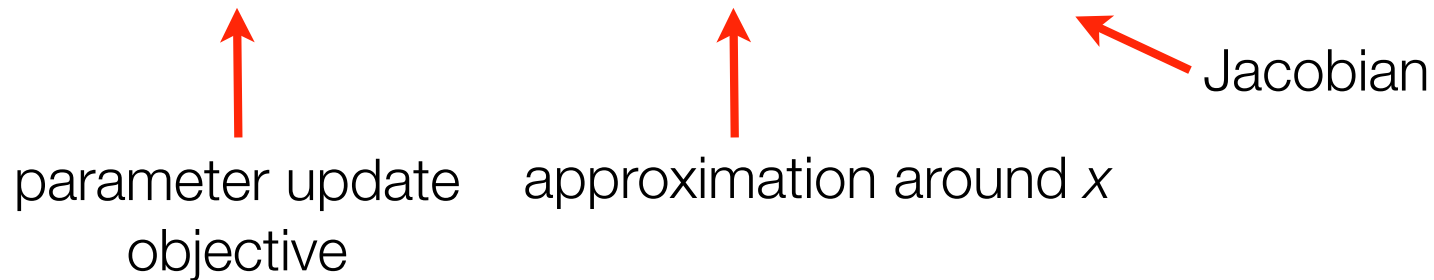
 approximation around x

 Jacobian

Gauss-Newton method

- Iteratively find parameter updates $\Delta \mathbf{x}$
- Perform a *first-order Taylor expansion* of the residual around the current \mathbf{x} :

$$\|f(\mathbf{x} - \Delta \mathbf{x}; \mathbf{A}) - \mathbf{b}\|^2 \approx \|f(\mathbf{x}; \mathbf{A}) + J(\mathbf{x})\Delta \mathbf{x} - \mathbf{b}\|^2$$


parameter update objective approximation around x Jacobian

- Note that the resulting residual approximation is linear in $\Delta \mathbf{x}$:
 - The parameter update $\Delta \mathbf{x}$ may be obtained via linear least squares

Gauss-Newton method

- Writing down the linear least-squares solution for $\Delta \mathbf{x}$, we obtain:

$$\Delta \mathbf{x} = \left(J(\mathbf{x})^\top J(\mathbf{x}) \right)^{-1} J(\mathbf{x})^\top r(\mathbf{x})$$



“Gauss-Newton approximation to Hessian”

- Gauss-Newton iteratively performs this update: $\mathbf{x} \leftarrow \mathbf{x} - \Delta \mathbf{x}$
- The Taylor expansion just became inaccurate! So iterate the whole process...

Gauss-Newton method

- Writing down the linear least-squares solution for $\Delta \mathbf{x}$, we obtain:

$$\Delta \mathbf{x} = \left(J(\mathbf{x})^\top J(\mathbf{x}) \right)^{-1} J(\mathbf{x})^\top r(\mathbf{x})$$



“Gauss-Newton approximation to Hessian”

- Gauss-Newton iteratively performs this update: $\mathbf{x} \leftarrow \mathbf{x} - \Delta \mathbf{x}$
- The Taylor expansion just became inaccurate! So iterate the whole process...

- To implement, you only need to derive *Jacobian*: $J(\mathbf{x}) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$

Newton's method

- Perform a *second-order Taylor expansion* of $g(\mathbf{x})$ around \mathbf{x} :

$$g(\mathbf{x}) \approx \|r(\mathbf{x})\|^2 - 2J(\mathbf{x})r(\mathbf{x})\Delta\mathbf{x} + [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^2] (\Delta\mathbf{x})^2$$

with residuals: $r(\mathbf{x}) = f(\mathbf{x}; \mathbf{A}) - \mathbf{b}$

Newton's method

- Perform a *second-order Taylor expansion* of $g(\mathbf{x})$ around \mathbf{x} :

$$g(\mathbf{x}) \approx \|r(\mathbf{x})\|^2 - 2J(\mathbf{x})r(\mathbf{x})\Delta\mathbf{x} + [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^2] (\Delta\mathbf{x})^2$$

with residuals: $r(\mathbf{x}) = f(\mathbf{x}; \mathbf{A}) - \mathbf{b}$

- This looks a lot like a linear least-squares problem; set gradient to zero:

$$-2J(\mathbf{x})^T r(\mathbf{x}) + 2 [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} = 0$$

$$[H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} = J(\mathbf{x})^T r(\mathbf{x})$$

Newton's method

- Perform a *second-order Taylor expansion* of $g(\mathbf{x})$ around \mathbf{x} :

$$g(\mathbf{x}) \approx \|r(\mathbf{x})\|^2 - 2J(\mathbf{x})r(\mathbf{x})\Delta\mathbf{x} + [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^2] (\Delta\mathbf{x})^2$$

with residuals: $r(\mathbf{x}) = f(\mathbf{x}; \mathbf{A}) - \mathbf{b}$

- This looks a lot like a linear least-squares problem; set gradient to zero:

$$\begin{aligned} -2J(\mathbf{x})^T r(\mathbf{x}) + 2 [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} &= 0 \\ [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})] \Delta\mathbf{x} &= J(\mathbf{x})^T r(\mathbf{x}) \end{aligned}$$

- Note the similarity of the Newton update with the Gauss-Newton update:

$$\Delta\mathbf{x} = [H(\mathbf{x})r(\mathbf{x}) + J(\mathbf{x})^T J(\mathbf{x})]^{-1} J(\mathbf{x})^T r(\mathbf{x})$$

Homographies

Homogeneous coordinates

- We are used to describing a location in *Cartesian coordinates*:

$$\mathbf{x} = [x \ y]^T$$

$$\mathbf{x} = [x \ y \ z]^T$$

Homogeneous coordinates

- We are used to describing a location in *Cartesian coordinates*:

$$\mathbf{x} = [x \ y]^T$$

$$\mathbf{x} = [x \ y \ z]^T$$

- Alternatively, we can describe locations in *homogeneous coordinates*:

$$\tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{w}]^T$$

$$\tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{z} \ \tilde{w}]^T$$

Homogeneous coordinates

- We are used to describing a location in *Cartesian coordinates*:

$$\mathbf{x} = [x \ y]^T \qquad \mathbf{x} = [x \ y \ z]^T$$

- Alternatively, we can describe locations in *homogeneous coordinates*:

$$\tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{w}]^T \qquad \tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{z} \ \tilde{w}]^T$$

- The corresponding Cartesian coordinates are given by:

$$\mathbf{x} = [\tilde{x}/\tilde{w} \ \tilde{y}/\tilde{w}]^T \qquad \mathbf{x} = [\tilde{x}/\tilde{w} \ \tilde{y}/\tilde{w} \ \tilde{z}/\tilde{w}]^T$$

Homogeneous coordinates

- We are used to describing a location in *Cartesian coordinates*:

$$\mathbf{x} = [x \ y]^T \qquad \mathbf{x} = [x \ y \ z]^T$$

- Alternatively, we can describe locations in *homogeneous coordinates*:

$$\tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{w}]^T \qquad \tilde{\mathbf{x}} = [\tilde{x} \ \tilde{y} \ \tilde{z} \ \tilde{w}]^T$$


- The corresponding Cartesian coordinates are given by:

$$\mathbf{x} = [\tilde{x}/\tilde{w} \ \tilde{y}/\tilde{w}]^T \qquad \mathbf{x} = [\tilde{x}/\tilde{w} \ \tilde{y}/\tilde{w} \ \tilde{z}/\tilde{w}]^T$$

- Essentially, you can think of \tilde{w} as a way to deal with object scale (“*disparity*”)
- Homogeneous coordinates are very useful when working with *perspective transformations (homographies)*

Pinhole camera

- *Pinhole camera* is a model for how an ideal camera works:



Pinhole camera


- *Focal length* of the camera influences what is captured on the image plane:



- A large focal length implies small *field of view*, and vice versa


How does a camera see a 3D point?

- Suppose we observe a 2D point, what is the corresponding *3D ray*?



How does a camera see a 3D point?

- Suppose we observe a 2D point, what is the corresponding 3D ray?



$$\mathbf{p} = [\mathbf{R}_s | \mathbf{c}_s] \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_s \\ y_s \\ 1 \end{bmatrix} = \mathbf{M}_s \bar{\mathbf{x}}_s$$

3D orientation
of image plane
(relative to nodal point)

3D location
of image plane
(relative to nodal point)


pixel spacing

2D pixel
(augmented notation)

How does a camera see a 3D point?

- The reverse operation has an unknown scaling (since we don't know the depth)
- For a *camera-centered* point, the 3D-to-2D projection can therefore be written in terms of the *calibration matrix* (note the homogeneous coordinates!):

$$\tilde{\mathbf{x}}_s = \alpha \mathbf{M}_s^{-1} \mathbf{p}_c = \mathbf{K} \mathbf{p}_c \approx \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \mathbf{p}_c$$



How does a camera see a 3D point?

- The reverse operation has an unknown scaling (since we don't know the depth)
- For a *camera-centered* point, the 3D-to-2D projection can therefore be written in terms of the *calibration matrix* (note the homogeneous coordinates!):

$$\tilde{\mathbf{x}}_s = \alpha \mathbf{M}_s^{-1} \mathbf{p}_c = \mathbf{K} \mathbf{p}_c \approx \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \mathbf{p}_c$$

- Including 3D rotation and location of the camera, we obtain:

$$\tilde{\mathbf{x}} = \mathbf{K} [\mathbf{R} | \mathbf{t}] \mathbf{p} = \mathbf{P} \mathbf{p}$$

camera intrinsics
(calibration matrix)

camera extrinsics
(rotation + translation)

camera matrix

How does a camera see a 3D point?

- The camera matrix is an example of a *projective transformation*:


166 pixels tall



370 pixels tall



600 pixels tall



How do two cameras see a place?

- Assume we have two cameras with projection matrices $\tilde{\mathbf{P}}_0$ and $\tilde{\mathbf{P}}_1$
- Where does camera 1 see the point that camera 0 sees at $\tilde{\mathbf{x}}_0$?

$$\tilde{\mathbf{x}}_1 \sim \tilde{\mathbf{P}}_1 \tilde{\mathbf{P}}_0^{-1} \tilde{\mathbf{x}}_0 = \mathbf{M}_{10} \tilde{\mathbf{x}}_0$$

How do two cameras see a place?

- Assume we have two cameras with projection matrices $\tilde{\mathbf{P}}_0$ and $\tilde{\mathbf{P}}_1$
- Where does camera 1 see the point that camera 0 sees at $\tilde{\mathbf{x}}_0$?

$$\tilde{\mathbf{x}}_1 \sim \tilde{\mathbf{P}}_1 \tilde{\mathbf{P}}_0^{-1} \tilde{\mathbf{x}}_0 = \mathbf{M}_{10} \tilde{\mathbf{x}}_0$$


- *Disparity* of point is irrelevant, so we can take the 3x3 sub-matrix of \mathbf{M}_{10} :

$$\tilde{\mathbf{x}}_1 \sim \tilde{\mathbf{H}}_{10} \tilde{\mathbf{x}}_0$$

- This is known as a *homography* $\tilde{\mathbf{H}}_{10}$ between the two cameras


How do two cameras see a place?

- Illustration of homography between two camera (for point and plane):



How do two cameras see a place?

- Illustration of homography between two camera (for point and plane):



$$H_{10} = \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & 1 \end{bmatrix}$$

- In Cartesian coordinates, the homography is given by:

$$x_1 = \frac{h_{00}x_0 + h_{01}y_0 + h_{02}}{h_{20}x_0 + h_{21}y_0 + 1}$$

$$y_1 = \frac{h_{10}x_0 + h_{11}y_0 + h_{12}}{h_{20}x_0 + h_{21}y_0 + 1}$$


Image stitching

Image stitching


- Stitching algorithms typically have four main ingredients:
 - Method to determine *correspondences* between images
 - Model describing the set of possible *motions* between images (homography)
 - Algorithm to perform *alignment* of the images (bundle adjustment)
 - Algorithm that *composites* the images after alignment (blending; seam finding)

Determining correspondences


- Feature detection:



- Feature description:




- Feature matching:



Motion model

- We will consider *homographies* because of their generality:



Fitting a homography

- Recall the definition of a homography in Cartesian coordinates:

$$x' = f(x, y) = \frac{h_{00}x + h_{01}y + h_{02}}{h_{20}x + h_{21}y + 1} \quad y' = g(x, y) = \frac{h_{10}x + h_{11}y + h_{12}}{h_{20}x + h_{21}y + 1}$$

- This makes the alignment objective a non-linear least squares problem:

$$\sum_i \left\| \begin{bmatrix} f(x_i, y_i; \mathbf{H}) \\ g(x_i, y_i; \mathbf{H}) \end{bmatrix} - \begin{bmatrix} x'_i \\ y'_i \end{bmatrix} \right\|^2$$

Fitting a homography

- Recall the definition of a homography in Cartesian coordinates:

$$x' = f(x, y) = \frac{h_{00}x + h_{01}y + h_{02}}{h_{20}x + h_{21}y + 1} \quad y' = g(x, y) = \frac{h_{10}x + h_{11}y + h_{12}}{h_{20}x + h_{21}y + 1}$$

- This makes the alignment objective a non-linear least squares problem:

$$\sum_i \left\| \begin{bmatrix} f(x_i, y_i; \mathbf{H}) \\ g(x_i, y_i; \mathbf{H}) \end{bmatrix} - \begin{bmatrix} x'_i \\ y'_i \end{bmatrix} \right\|^2$$

- Simply use Gauss-Newton's (or Newton's) method to solve this problem:
 - You may have to use RANSAC to deal with outliers!

Remark on RANSAC

- The treatment of RANSAC we saw last week was somewhat simplified
- It was not random at all!

Remark on RANSAC

- The treatment of RANSAC we saw last week was somewhat simplified
- It was not random at all! Full RANSAC actually works as follows:
 - 1) Select random subset from 50% of “best” matches as initial inliers
 - 2) Model is fitted to the *hypothetical inliers*
 - 3) Data are tested against the fitted model to determine hypothetical inliers
 - 4) Return to step 2) until sufficient points are classified as inliers (or fixed number of times)
 - 5) Return to step 1) a fixed number of times


Fitting a homography

- Consider the following two images and SIFT matches between them:




Fitting a homography

- Consider the following two images and SIFT matches between them:



Fitting a homography

- Visualization of the inliers found by RANSAC:




Fitting a homography

- Visualization of the homography found between the two images:




Alpha blending

- *Alpha blending averages the images; leads to “ghosting” and visible seams:*



Finding an optimal seam


- We would like to find a *seam* that minimizes the difference between images
- This can be formulated as a *graph min-cut* problem, as follows:



- When you want pixel \mathbf{x} to be taken from I_1 , set $f_2(\mathbf{x}) = \infty$ (and vice versa)
- In all other cases: $f_1(\mathbf{x}) = f_2(\mathbf{x}) = 0$
- Set, e.g., $g(x, y) = |I_1(\mathbf{x}) - I_2(\mathbf{y})| + |I_2(\mathbf{x}) - I_1(\mathbf{y})|$
- Efficient graph-cut algorithms exist


Finding an optimal seam

- Ghosting can be prevented by finding a *seam* at which to switch images:




Feathering

- Visibility of seam may be reduced by “*blurring*” weights of both images:




Correcting for exposure differences

- Instead of using seam to select image, we use it to select the *image gradient*:






Correcting for exposure differences

- The stitched image can then be recovered from this gradient:




Lens distortion

- Homography model may be too simple when lens distortion is present:



- A common way of dealing with this is by incorporating a distortion model in the objective, and also minimizing w.r.t. the parameters of that model

Other applications of homographies



Reading material: Section 2 and 9