

Face Processing

IN4393 – Computer Vision

Introduction

- Face processing/analysis comprises a number of different tasks:
 - Face detection (*“where is a face?”*)
 - Face recognition (*“of whom is this face?”*)
 - Face verification (*“are these faces the same?”*)
 - Expression recognition (*“is this face happy or not?”*)

Face detection: Viola & Jones

- Train a *classifier* to predict whether a bounding box contains a face or not:



Face detection: Viola & Jones


- At test time, use a sliding window detector for multiple scales:



- Use *non-maxima suppression* in x-y-scale space to filter classifier predictions

Face detection: Viola & Jones

- Extract *Haar features* from the image patch, using the *integral image*:



$$\text{Feature Value} = \sum (\text{pixels in white area}) - \sum (\text{pixels in black area})$$



- We find features that are common in faces, and use these as *weak learners*

Face detection: Viola & Jones

- Features can be computed efficiently using the *integral image*:

12	8	2	4	7		12	20	22	26	33
2	11	3	6	8		14	33	38	48	63
3	2	0	1	10	→	17	38	43	54	79
1	5	2	7	2		18	44	51	69	96
0	0	2	3	2		18	44	53	74	103

Face detection: Viola & Jones

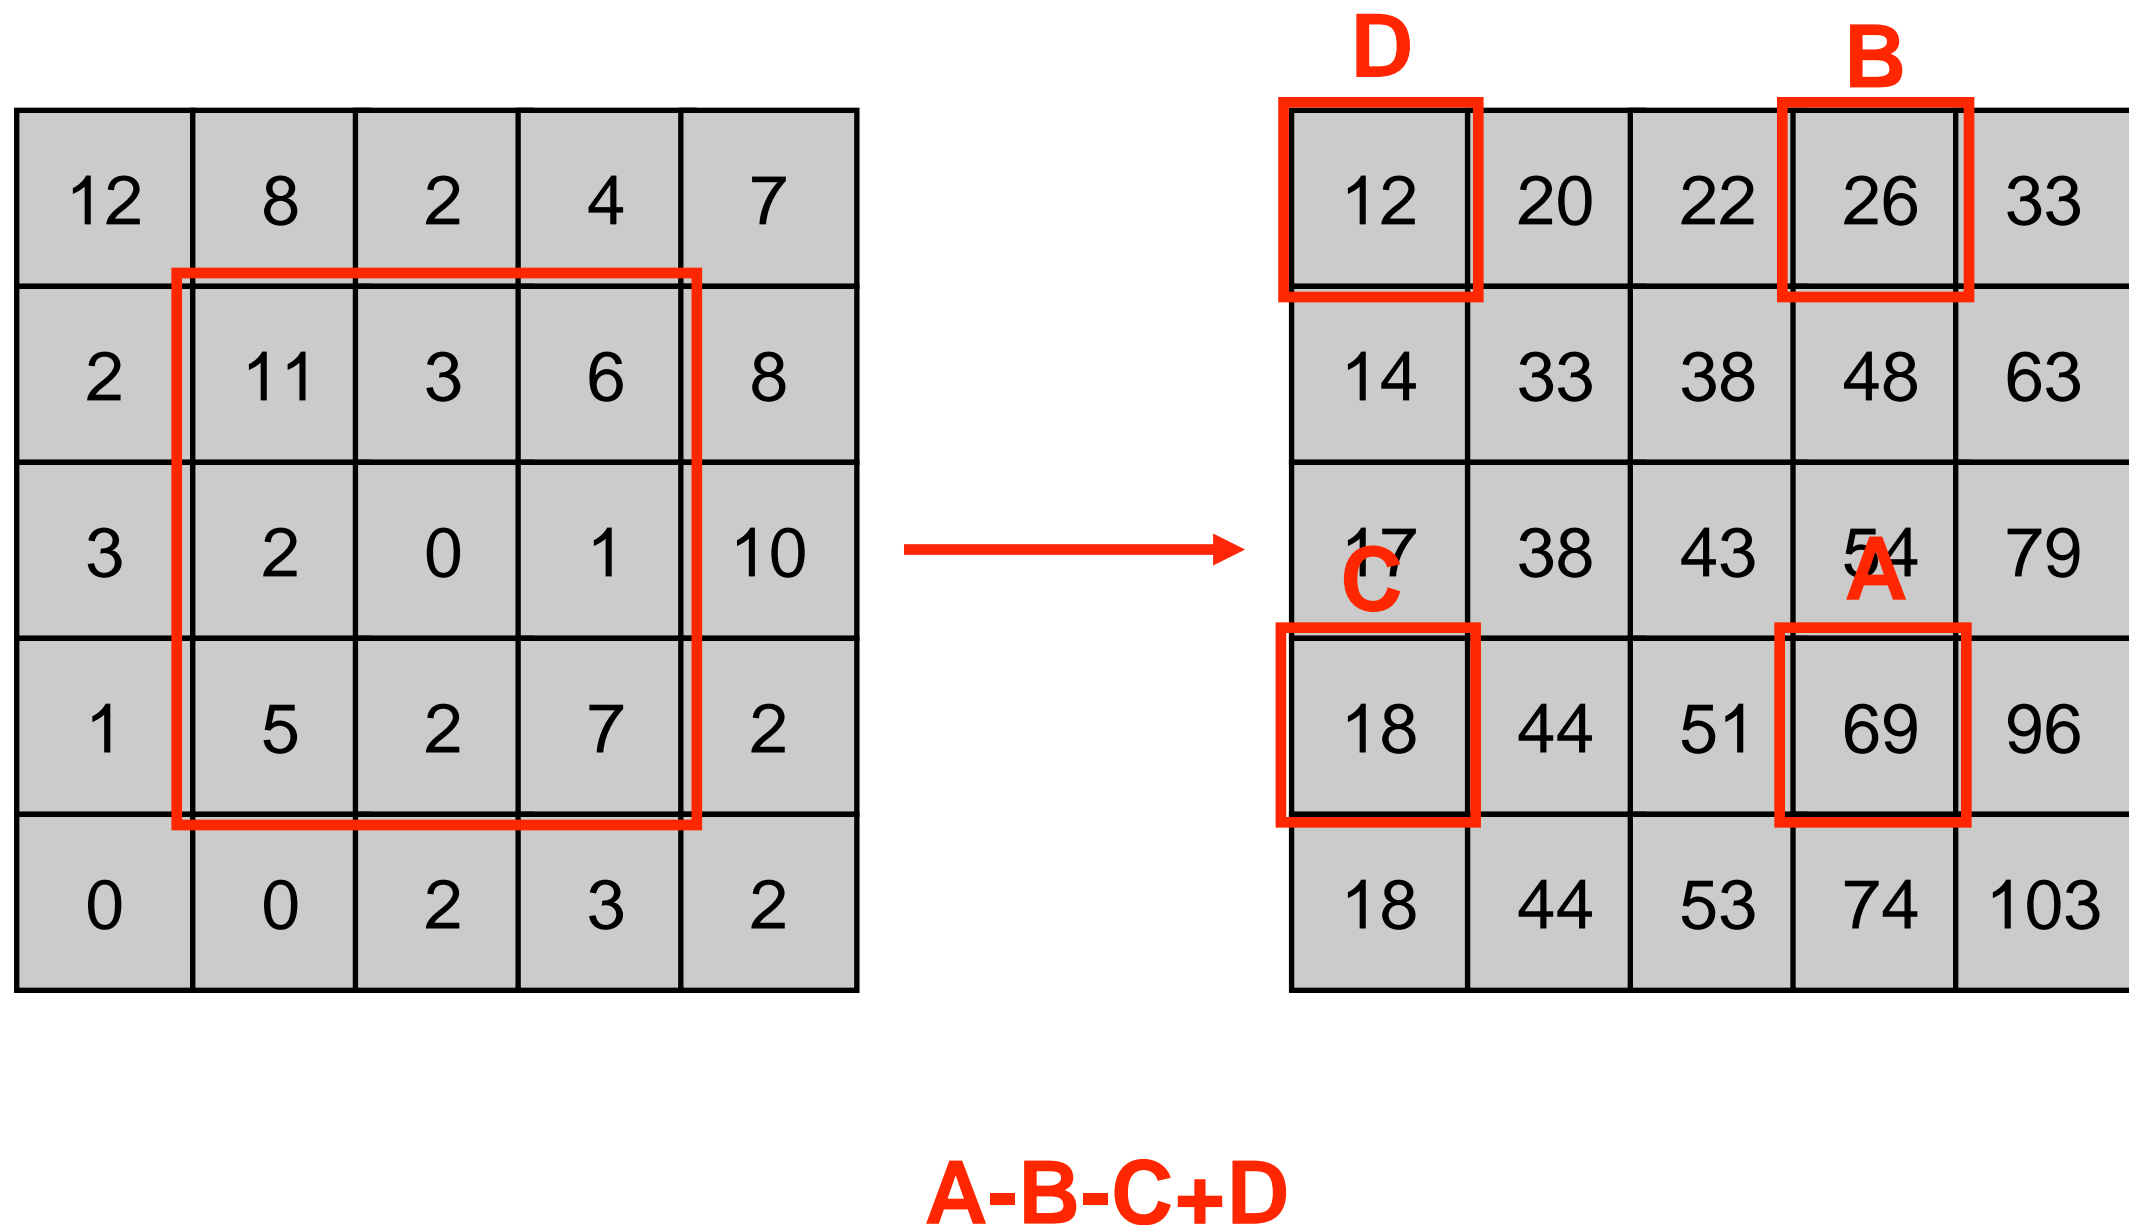
- Features can be computed efficiently using the *integral image*:

12	8	2	4	7
2	11	3	6	8
3	2	0	1	10
1	5	2	7	2
0	0	2	3	2

12	20	22	26	33
14	33	38	48	63
17	38	43	54	79
18	44	51	69	96
18	44	53	74	103

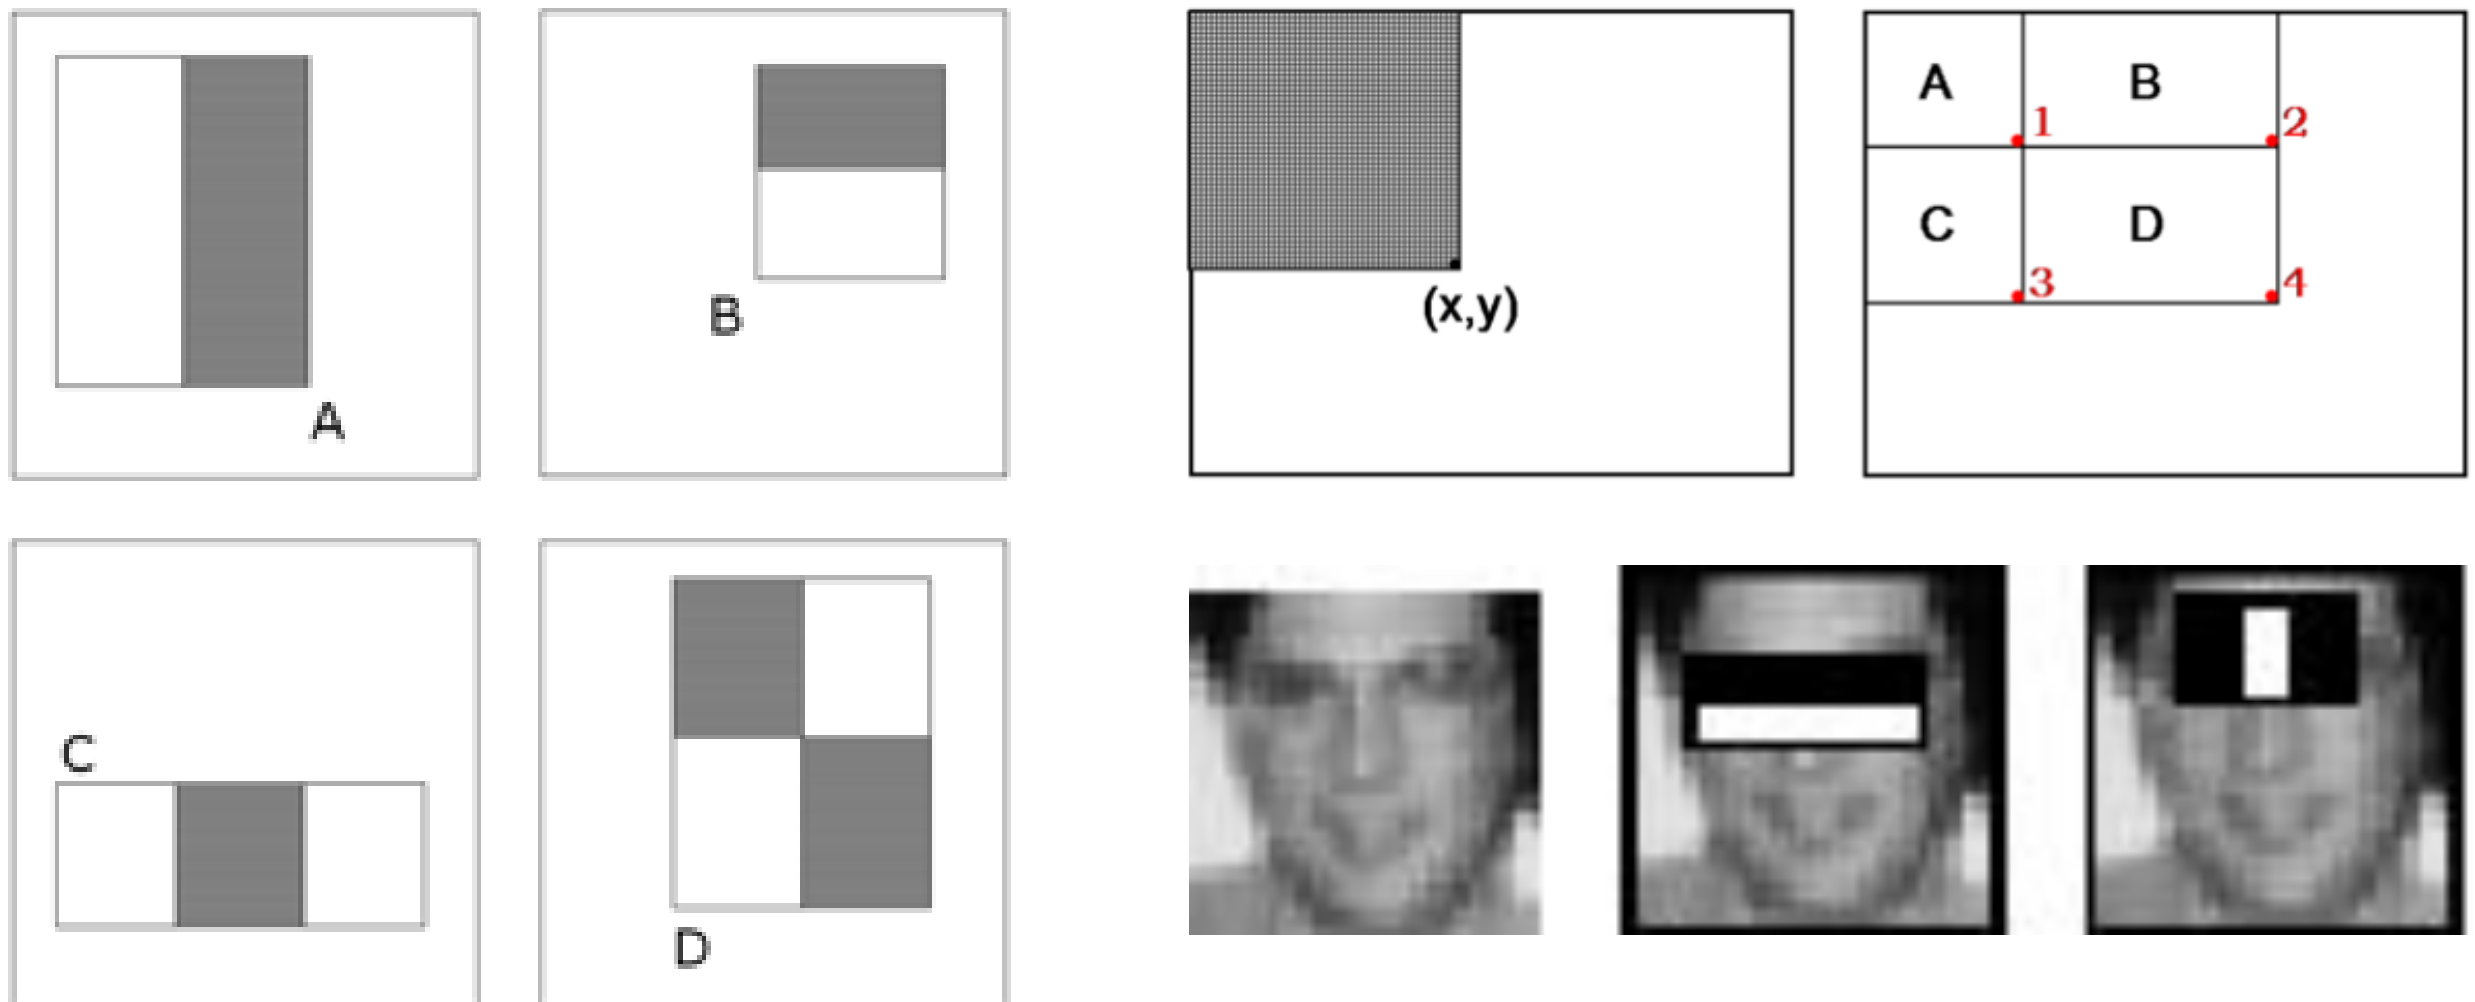
Face detection: Viola & Jones

- Features can be computed efficiently using the *integral image*:



Face detection: Viola & Jones

- Extract *Haar features* from the image patch, using the *integral image*:



- We find features that are common in faces, and use these as *weak learners*

Face detection: Viola & Jones

- *AdaBoost* training on annotated data set; learns collection of *weak learners*:

$$h(\mathbf{x}) = \text{sign} \left[\sum_{i=1}^m \alpha_i h_i(\mathbf{x}) \right]$$

Face detection: Viola & Jones

- *AdaBoost* training on annotated data set; learns collection of *weak learners*:

$$h(\mathbf{x}) = \text{sign} \left[\sum_{i=1}^m \alpha_i h_i(\mathbf{x}) \right]$$

- In Viola & Jones, the weak learners are *decision stumps*:

$$h_i(\mathbf{x}) = [f_i \geq \theta_i]$$

Face detection: Viola & Jones

- *AdaBoost* training on annotated data set; learns collection of *weak learners*:

$$h(\mathbf{x}) = \text{sign} \left[\sum_{i=1}^m \alpha_i h_i(\mathbf{x}) \right]$$

- In Viola & Jones, the weak learners are *decision stumps*:


$$h_i(\mathbf{x}) = [f_i \geq \theta_i]$$

- Iteratively select the learner that minimizes the *weighted* classification error:


$$e_i = \sum_{n=1}^N w_{n,i} (1 - \delta(y_n, h_i(\mathbf{x}_n; \theta_i)))$$

- Efficient algorithms exist to find the threshold in linear time
- Update the *per-instance weights* based on the classification error of weak learner


AdaBoost learning illustration




AdaBoost learning illustration




AdaBoost learning illustration



AdaBoost learning illustration




AdaBoost learning illustration




AdaBoost learning illustration

**Final classifier is
a combination of
weak classifiers**



Face detection

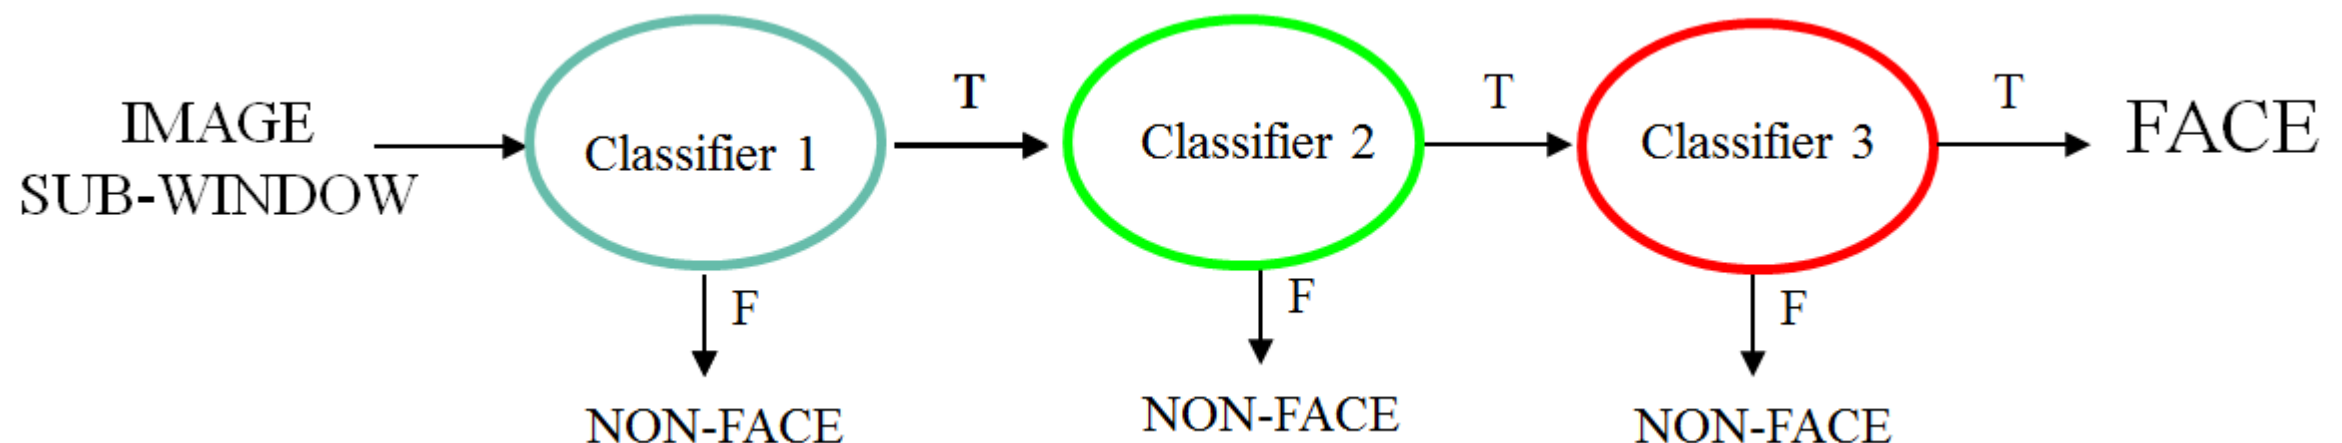
- Schematic overview of AdaBoost learning with decision stumps:



- Update for per-instance weights:
$$w_{n,i+1} \leftarrow w_{n,i} \left(\frac{e_i}{1 - e_i} \right)^{1 - \delta(y_n, h_i(\mathbf{x}_n; \theta_i))}$$
- Weak-learner weights given by:
$$\alpha_i = -\log \left(\frac{e_i}{1 - e_i} \right)$$

Face detection: Viola & Jones

- To perform the detection, we use a sliding window detector (at multiple scales)
- The classification of a patch can be performed using a *cascaded classifier*.



Note that this is extremely fast at test time: for negative examples, we typically only need to compute a very small number of features!

Face detection: Viola & Jones


- False positive rate of a cascade with K classifiers: $FPR = \prod_{i=1}^K FPR_i$
- Detection rate of a cascade with K classifiers: $DR = \prod_{i=1}^K DR_i$

Face detection: Viola & Jones

- False positive rate of a cascade with K classifiers: $FPR = \prod_{i=1}^K FPR_i$
- Detection rate of a cascade with K classifiers: $DR = \prod_{i=1}^K DR_i$
- Assume we have a cascade of $K = 32$ classifiers:
 - To get a false positive rate of 10^{-6} , each classifier may have FPR of 65%
 - To get a detection rate of 90%, each classifier should have DR of 99.7%




Face detection: Viola & Jones

- Multiple locations near a face will typically yields multiple detections
- In the original V&J detector, the detections are post-processed as follows:
 - Whenever two detections overlap, the bounding boxes are merged
 - The final detection is the average of the corners of all merged detections:



Face detection: Viola & Jones

- Examples of face detections (using V&J implementation in OpenCV):



Face detection: Viola & Jones

- Detection of profile faces requires training on separate data set:




Face recognition: Eigenfaces

- Consider face images as high-dimensional data points
- Apply *dimension reduction* on the images to obtain low-dimensional features
- The reduction is performed using *principal components analysis*

Face recognition: Eigenfaces

- Principal Components Analysis maps the data in a *linear subspace*, such that the *variance* of the projected data is maximized:



Face recognition: Eigenfaces


- Our objective is to maximize variance: $\max_{\|\mathbf{w}\|^2=1} \text{var}(\mathbf{w}^T \mathbf{X})$
- Assuming zero-mean data: $\text{var}(\mathbf{w}^T \mathbf{X}) = [\mathbf{w}^T \mathbf{X} \mathbf{X}^T \mathbf{w}] = [\mathbf{w}^T \mathbf{C} \mathbf{w}]$
- Enforce constraint using *Lagrange multipliers*:

$$\max_{\|\mathbf{w}\|^2=1} \text{var}(\mathbf{w}^T \mathbf{X}) = \max_{\mathbf{w}, \lambda} \mathbf{w}^T \mathbf{C} \mathbf{w} - \lambda(\mathbf{w}^T \mathbf{w} - 1)$$

- Set gradient with respect to \mathbf{w} to zero: $\mathbf{C} \mathbf{w} - \lambda \mathbf{w} = 0$
 $\mathbf{C} \mathbf{w} = \lambda \mathbf{w}$


Face recognition: Eigenfaces

- We can move through the PCA subspace to generate new faces:



Face recognition: Eigenfaces

- We can visualize the eigenfaces to show the main sources of variation:




- We may use $\mathbf{z} = [\mathbf{w}_1^T \mathbf{x}, \dots, \mathbf{w}_p^T \mathbf{x}]^T$ as features for identity recognition

Face recognition: Active appearance models

- Separates face variations into *shape* and *texture* variation:

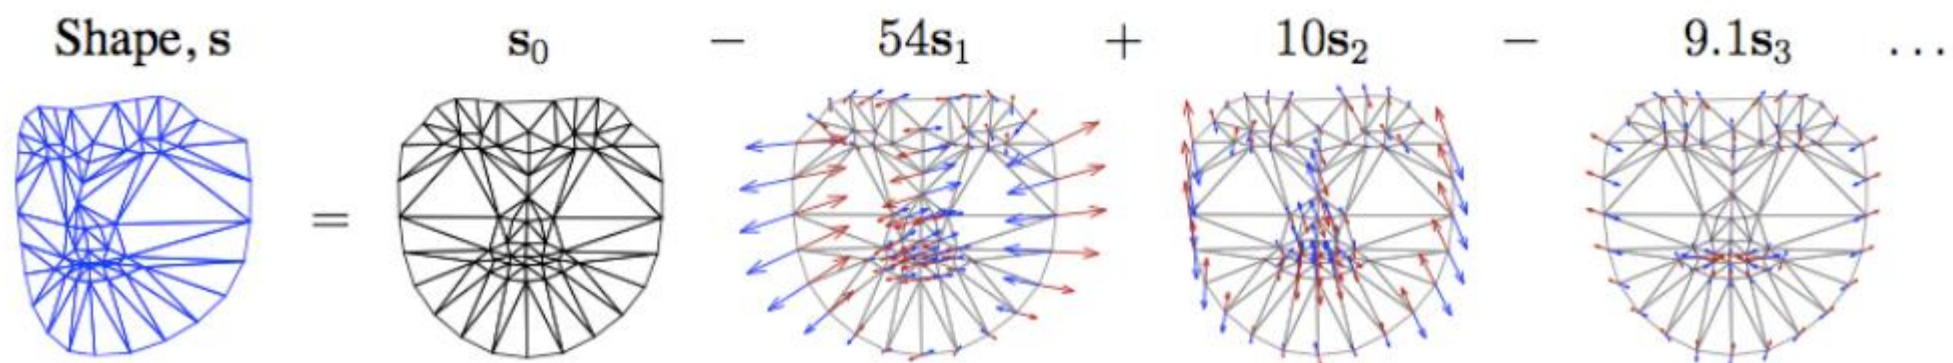


Appearance, A = A_0 + $3559A_1$ + $351A_2$ - $256A_3$...




AAM Model Instance
 $M(W(\mathbf{x}; \mathbf{p}))$

Shape, s = s_0 - $54s_1$ + $10s_2$ - $9.1s_3$...




Face recognition: Active appearance models

- Gather a data set of face images with annotated feature points
- Remove translations and rotations from point annotations (Procrustes alignment)
- Learn *point distribution model* and *texture model* from the data




Face recognition: Active appearance models

- *Point distribution model* is obtained using PCA:



- New facial shape is generated by linearly combining the components:

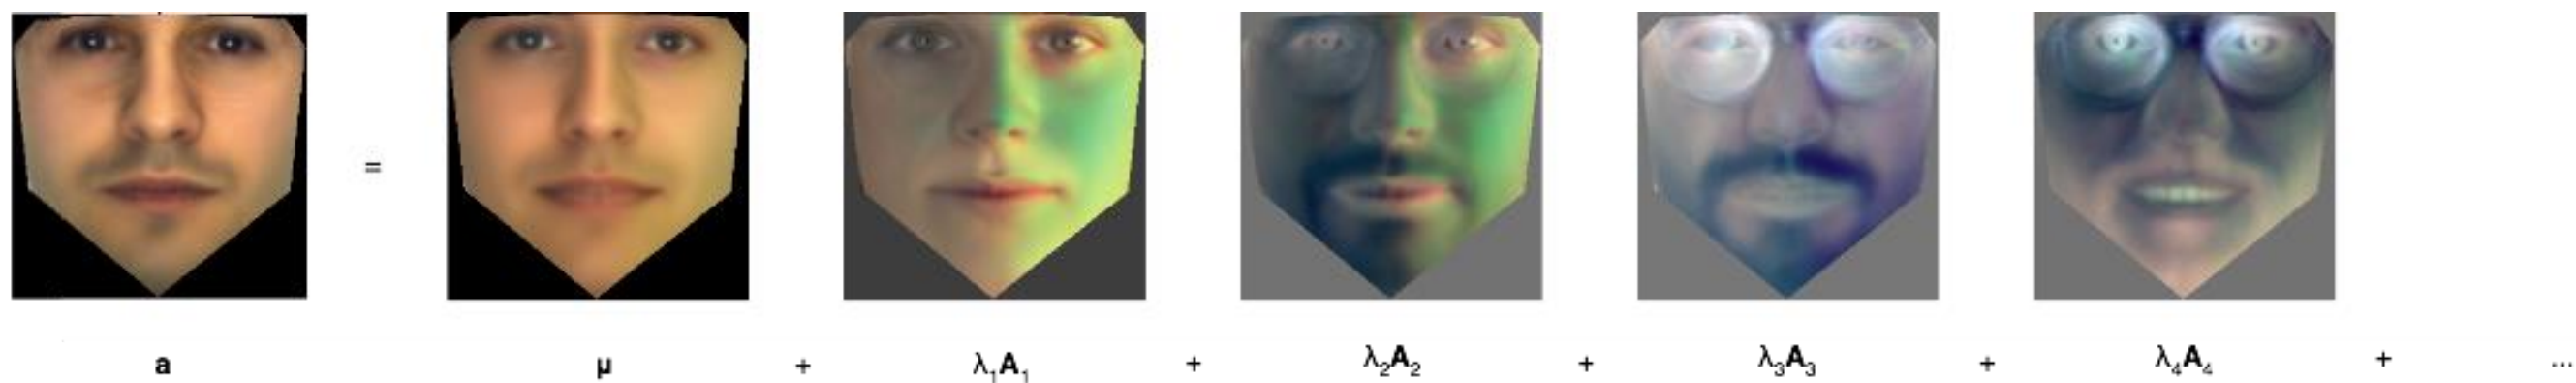


Face recognition: Active appearance models

- *Texture model* is also obtained using PCA:



- New facial textures are generated using a linear combination of components:



$\mathbf{a} = \boldsymbol{\mu} + \lambda_1 \mathbf{A}_1 + \lambda_2 \mathbf{A}_2 + \lambda_3 \mathbf{A}_3 + \lambda_4 \mathbf{A}_4 + \dots$

Face recognition: Active appearance models


- We receive a new face image in which we want to measure facial features:




- Now what do we do to fit the active appearance model to this new face?
 - Find parameters of the model that best fit the face (minimizing sum of squared errors) using *Lucas-Kanade algorithm*

Image warp

- We can *warp* between an arbitrary shape and the base shape (and vice versa):



$W(\mathbf{x}; \mathbf{p})$



$W(\mathbf{x}; \mathbf{p})$

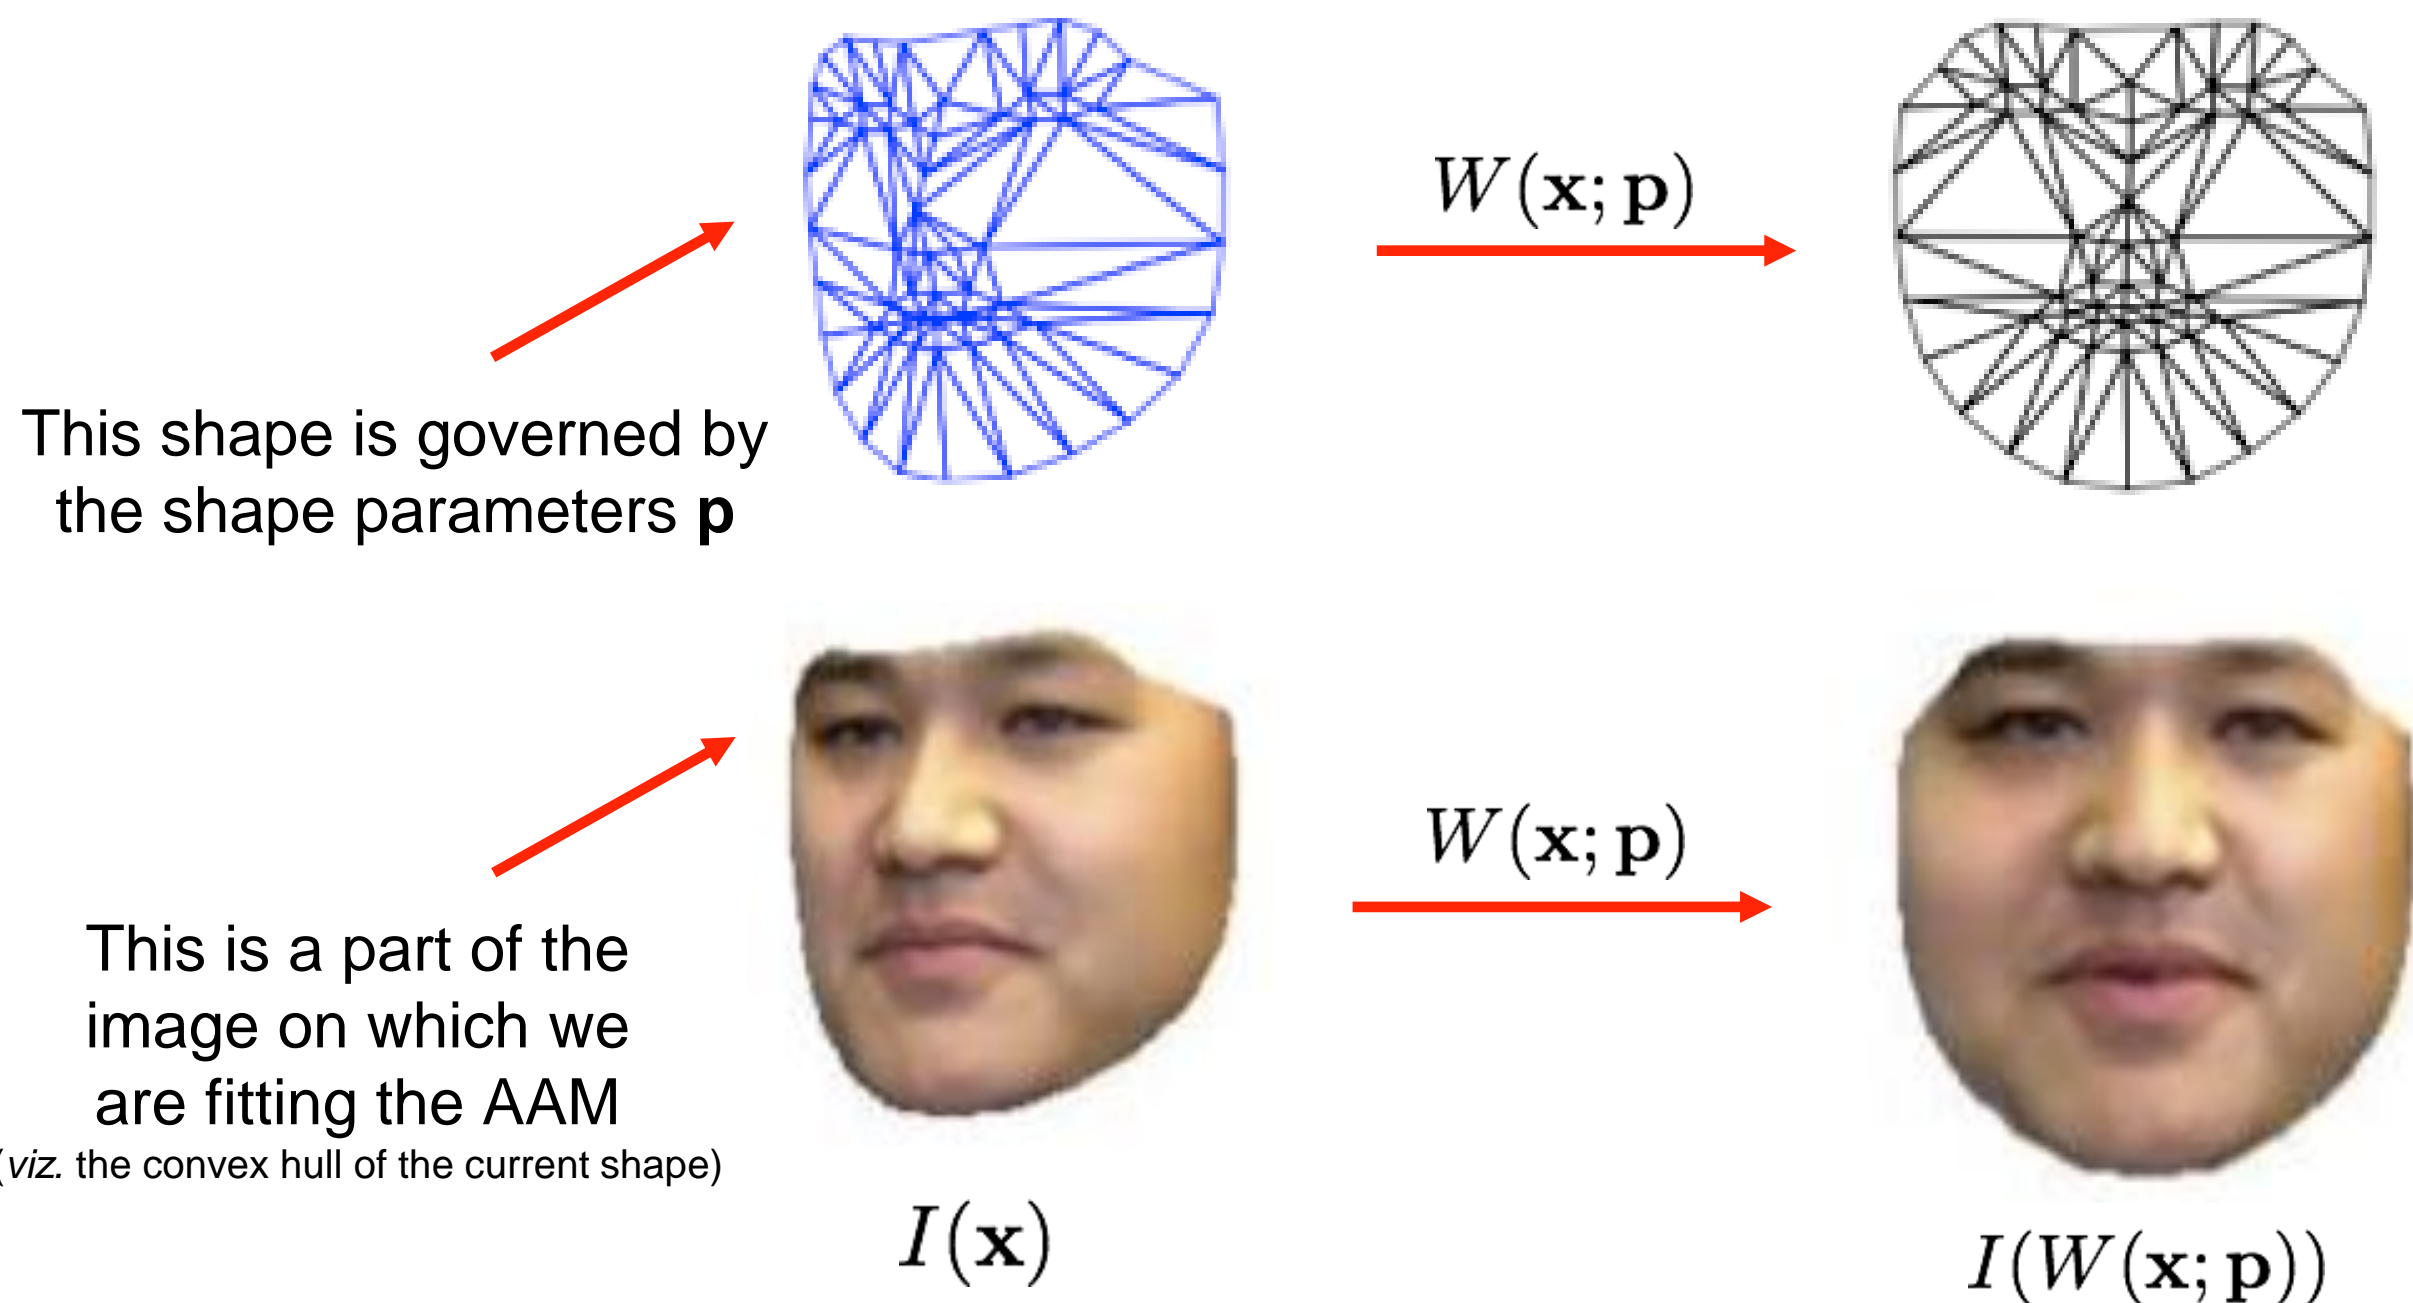


$I(\mathbf{x})$

$I(W(\mathbf{x}; \mathbf{p}))$

Image warp

- We can *warp* between an arbitrary shape and the base shape (and vice versa):



Lucas-Kanade algorithm

- Minimizes sum of squared error w.r.t. \mathbf{p} using *Gauss-Newton* algorithm:

$$\sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]^2$$



Mean texture

-



Test image
after warp

Lucas-Kanade algorithm

- Minimizes sum of squared error w.r.t. \mathbf{p} using *Gauss-Newton* algorithm:

$$\sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]^2$$



Mean texture

-



Test image
after warp

Goal: Set the shape parameters \mathbf{p} such that the left image looks as much as possible like the right image

Lucas-Kanade algorithm

- Minimizes sum of squared error w.r.t. \mathbf{p} using *Gauss-Newton* algorithm:

$$\sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]^2$$

- Iteratively solve for parameter increment $\Delta \mathbf{p}$:

$$\min_{\Delta \mathbf{p}} \sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p}))]^2$$

Lucas-Kanade algorithm

- Minimizes sum of squared error w.r.t. \mathbf{p} using *Gauss-Newton* algorithm:

$$\sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]^2$$

- Iteratively solve for parameter increment $\Delta \mathbf{p}$:

$$\min_{\Delta \mathbf{p}} \sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p}))]^2$$

- This is strongly non-linear, so write down *first-order Taylor expansion*:

Lucas-Kanade algorithm

- Minimizes sum of squared error w.r.t. \mathbf{p} using *Gauss-Newton* algorithm:

$$\sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]^2$$

- Iteratively solve for parameter increment $\Delta \mathbf{p}$:

$$\min_{\Delta \mathbf{p}} \sum_{\mathbf{x}} [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p}))]^2$$

- This is strongly non-linear, so write down *first-order Taylor expansion*:

$$\min_{\Delta \mathbf{p}} \sum_{\mathbf{x}} \left[A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p})) - \nabla I \frac{\partial W}{\partial \mathbf{p}} \Delta \mathbf{p} \right]^2$$

- As expected, this is a standard *linear least squares* problem

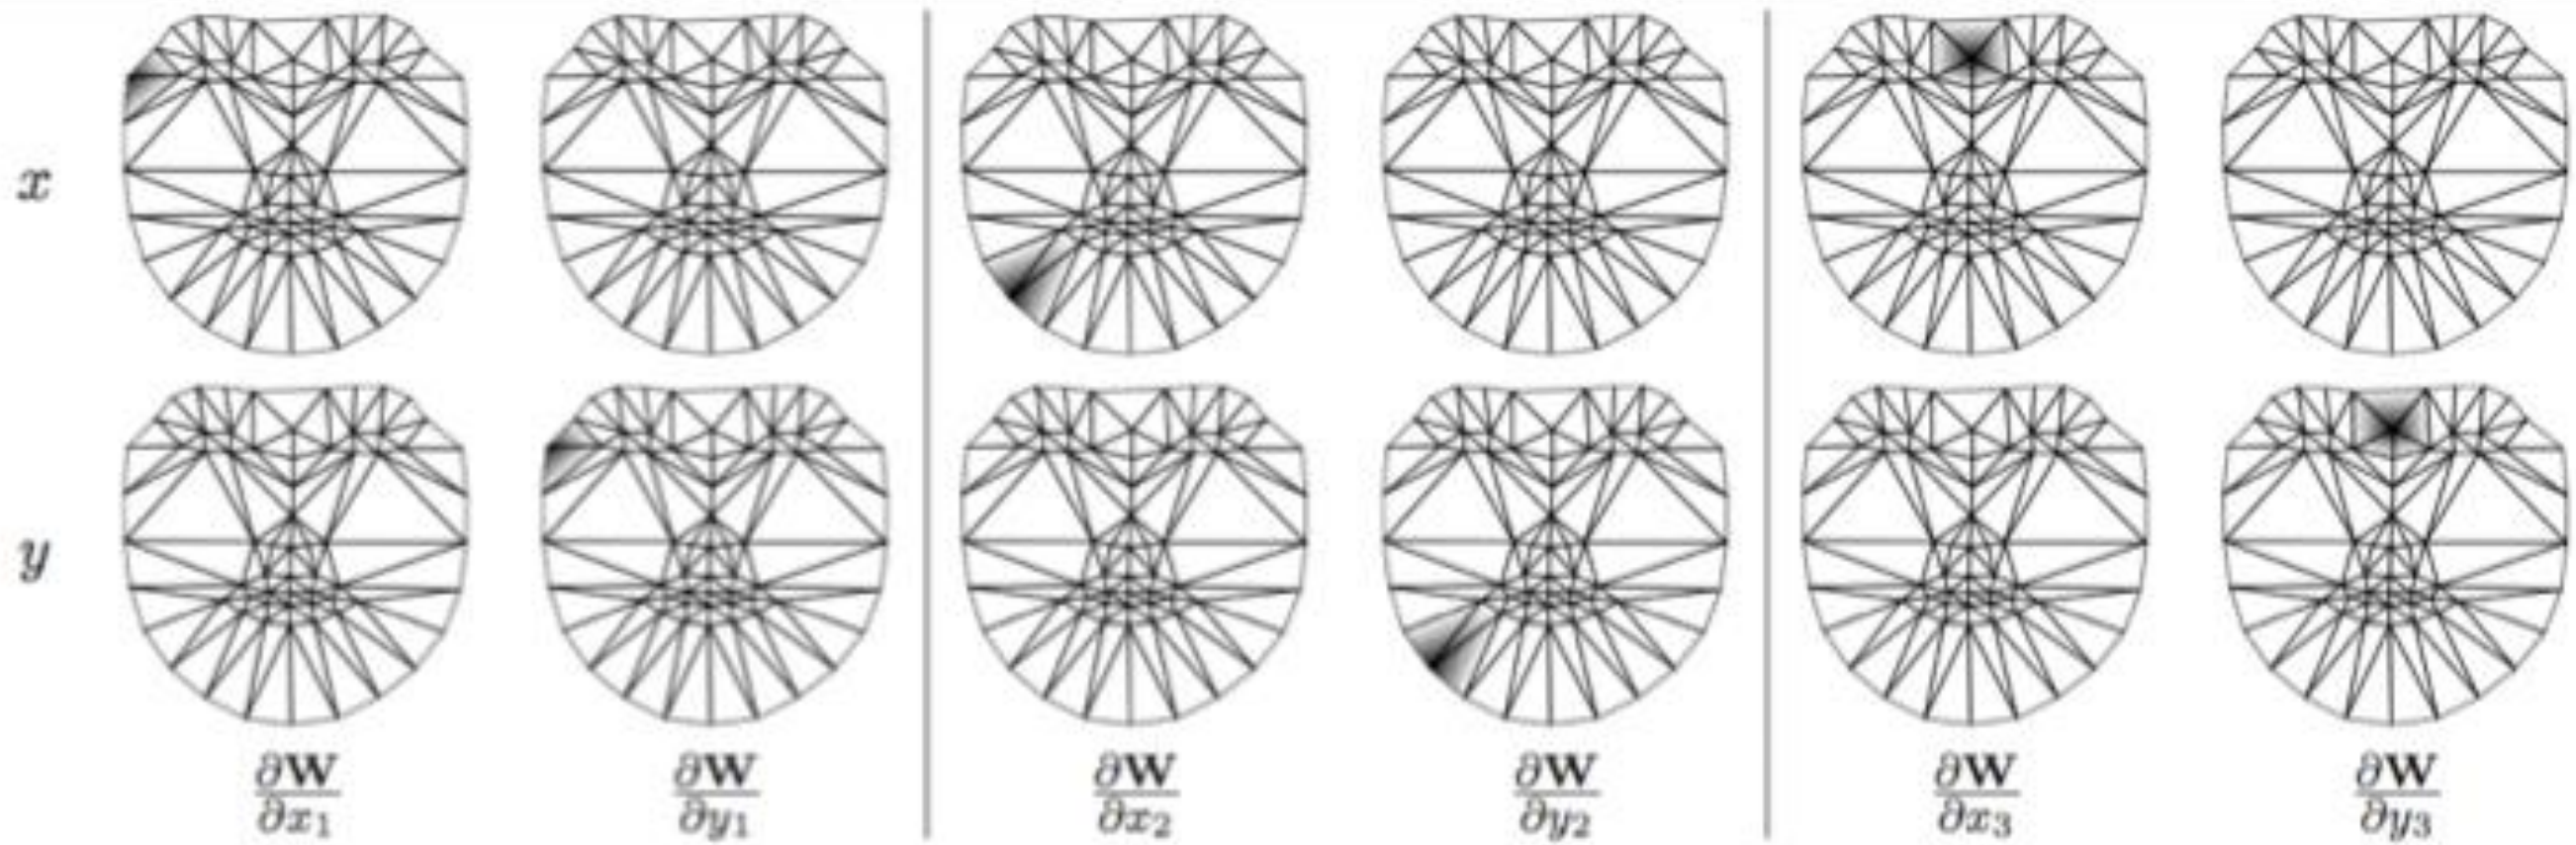
Gradient Images

- Illustration of the image gradient ∇I




Lucas-Kanade Algorithm

- Illustration of the warp Jacobian: $\frac{\partial W}{\partial \mathbf{p}} = \frac{\partial W}{\partial \mathbf{x}} \frac{\partial \mathbf{x}}{\partial \mathbf{p}}$



Lucas-Kanade Algorithm

- Illustration of the warp Jacobian: $\frac{\partial W}{\partial \mathbf{p}} = \frac{\partial W}{\partial \mathbf{x}} \frac{\partial \mathbf{x}}{\partial \mathbf{p}}$



Lucas-Kanade algorithm

- Closed-form solution for the parameter update:

$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial W}{\partial \mathbf{p}} \right]^T [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]$$

- Herein, \mathbf{H} is the Gauss-Newton approximation to the Hessian:

$$\mathbf{H} = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial W}{\partial \mathbf{p}} \right]^T \left[\nabla I \frac{\partial W}{\partial \mathbf{p}} \right]$$

Lucas-Kanade algorithm

- Closed-form solution for the parameter update:

$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_{\mathbf{x}} \left[\nabla I \frac{\partial W}{\partial \mathbf{p}} \right]^T [A_0(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]$$

- Herein, \mathbf{H} is the Gauss-Newton approximation to the Hessian:

$$\mathbf{H} = \sum_{\mathbf{x}} \left[\nabla I \frac{\partial W}{\partial \mathbf{p}} \right]^T \left[\nabla I \frac{\partial W}{\partial \mathbf{p}} \right]$$

- Every iteration requires computation of warp Jacobian and Hessian

Lucas-Kanade Algorithm

1. Warp I with $\mathbf{W}(\mathbf{x}; \mathbf{p}) \Rightarrow I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$
2. Compute error image $A_o(x) - I(\mathbf{W}(\mathbf{x}; \mathbf{p}))$
3. Warp gradient of I to compute ∇I
4. Evaluate Jacobian $\frac{\partial \mathbf{W}}{\partial \mathbf{p}}$
5. Compute Hessian
6. Compute $\Delta \mathbf{p}$
7. Update parameters $\mathbf{p} \leftarrow \mathbf{p} + \Delta \mathbf{p}$

Face recognition: Active appearance models



Face recognition: Active appearance models




Face recognition: Active appearance models



Face recognition: Active appearance models


- Illustration of fitting a shape model:



Face recognition: Active appearance models



Face recognition: Active appearance models



Face recognition and expression analysis

- Facial feature points (landmarks) can be used for a number of tasks:
 - *Facial expression analysis:*
 - Measure variations of landmark locations over time (shape variation); use texture features to measure presence of wrinkles (texture variation), *etc.*

Face recognition and expression analysis

- Facial feature points (landmarks) can be used for a number of tasks:
 - *Facial expression analysis:*
 - Measure variations of landmark locations over time (shape variation); use texture features to measure presence of wrinkles (texture variation), *etc.*
 - *Facial identity recognition or face verification* (passport control):
 - Measure characteristics that are invariant under expressions but person-specific: inter-ocular distance, relative position of nose, *etc.*
 - Build skin models:



Example: Recognition of Action Units (FACS)



AU 1
Inner brow raise



AU 2
Outer brow raise



AU 4
Brow lower



AU 6
Cheek raise



AU 9
Nose wrinkler



AU 12
Lip corner pull



AU 15
Lip corner depress




AU 20
Lip stretcher

Example: Recognition of Action Units (FACS)

- Inner brow raiser:
- Outer brow raiser:
- Brow lowerer:
- Upper lid raiser:
- Nose wrinkler:
- Lip corner depressor:
- Etcetera...



Example: Recognition of Action Units (FACS)



Example: Expression Recognition



Neutral	0 %
Happy	0 %
Surprise	0 %
Angry	0 %
Disgust	0 %
Fear	0 %
Sad	0 %


Status:

- * Source: Webcam
- * Player: Playing
- * Face: Found
- * Markers: Scale to face


Hint:

Press "Fit mask" to fit the

Example: Detection of Expression Spontaneity



Example: Detection of Expression Spontaneity



Example: Expression cloning




Example: Expression cloning



MSc Projects

- Are you interested in face processing?
 - Face Tracking
 - Expression Analysis
 - Age Estimation
 - Face Recognition
 - Kinship Verification/Recognition
 - Automatic Assessment of Depression
 - And more... (only limit is your imagination)



Reading material:

Section 14.1 and 14.2

Section 1 and 2 of “*Lucas-Kanade 20 Years on*”